

AZURE AKS Create Cluster Introduction



AKS - Introduction

- AKS - Azure Kubernetes Service
- AKS is highly available, secure and fully managed Service
- When compared to other cloud providers, AKS is the one which is available in highest number of regions
- Will be able to run any type of workloads
 - Windows based applications like .Net Apps
 - Linux supported applications like Java
 - IOT device deployment and management on demand
 - Machine Learning Model training with AKS
- Able to run in Hybrid Platforms
 - Azure Stack HCI
 - Windows Servers with Linux Distros
 - Planing for Vmware Platform
- Able to use Azure services without additional infra and admin effort
 - You can deploy and integrate azure services with your AKS easily
 - Azure Storage, Azure Key Vault, Azure Devops, Azure LB, etc.

MASTER

AKS Kube
Controller
Manager

kube-apiserver

etcd

kube-
scheduler

Container Runtime (Docker)

AKS cluster Control Plane

Worker Node 1

Kubelet

Kube-Proxy

Container Runtime (Docker)

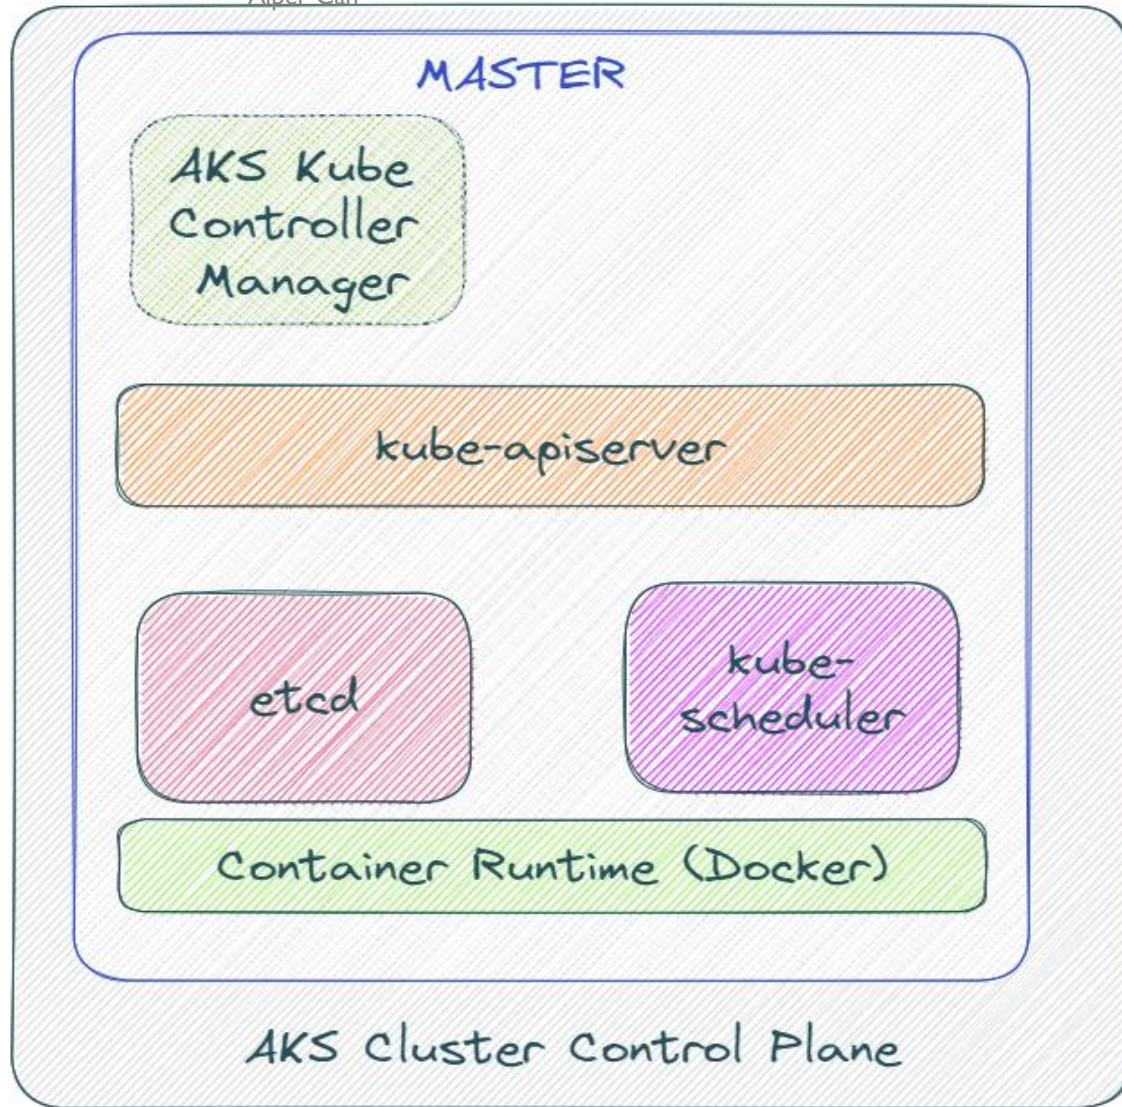
Worker Node 2

Kubelet

Kube-Proxy

Container Runtime (Docker)

AKS Node Pool



○ kube-apiserver

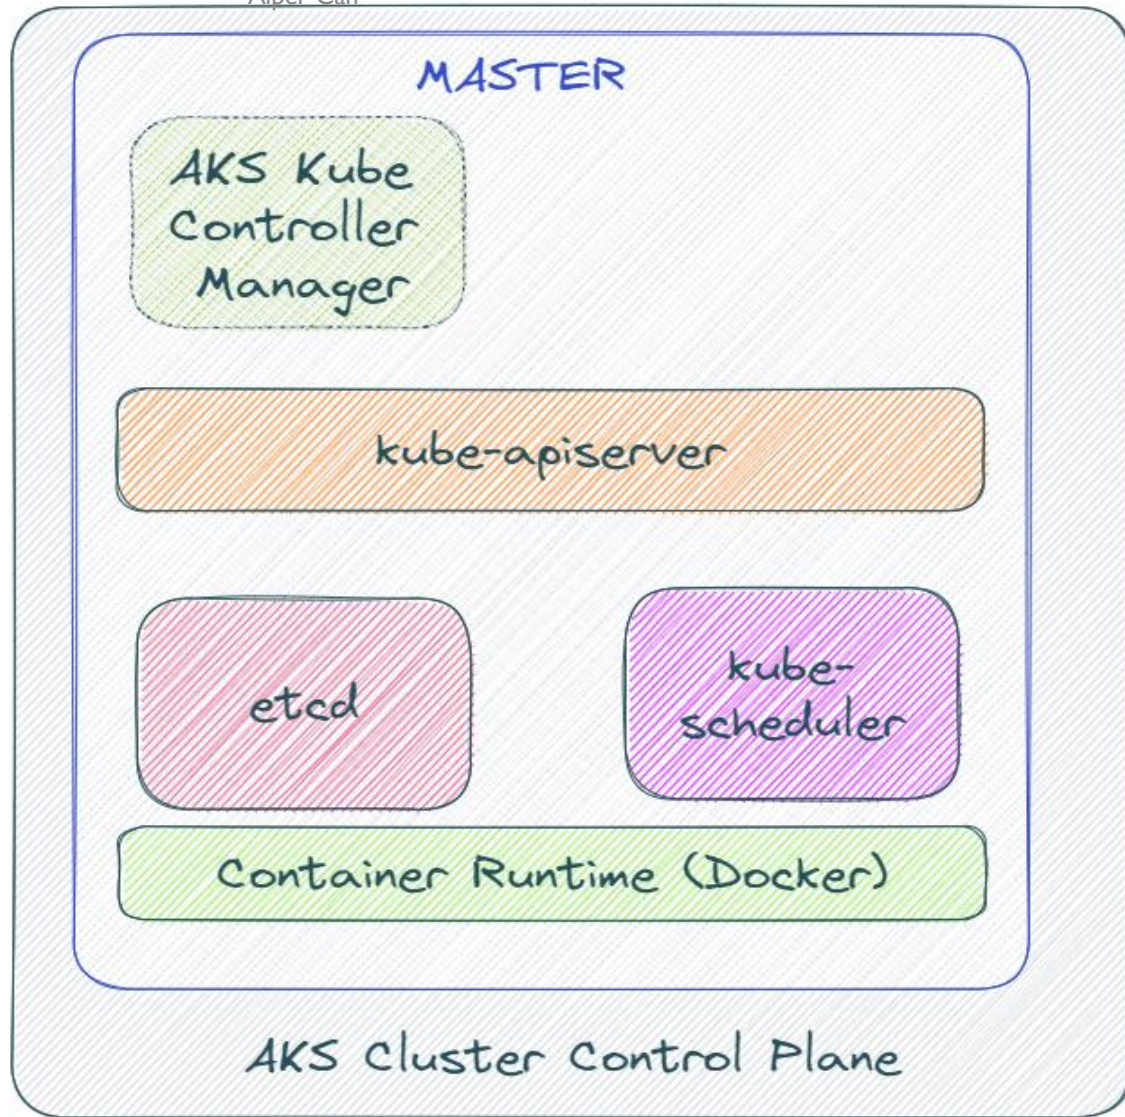
- It acts as front end for the Kubernetes control plane. It exposes the Kubernetes API
- Command line tools (like kubectl), Users and even Master components (scheduler, controller manager, etcd) and Worker node components like (Kubelet) everything talk with API Server.

○ etcd

- Consistent and highly-available key value store used as Kubernetes' backing store for all cluster data.
- It stores all the masters and worker node information.

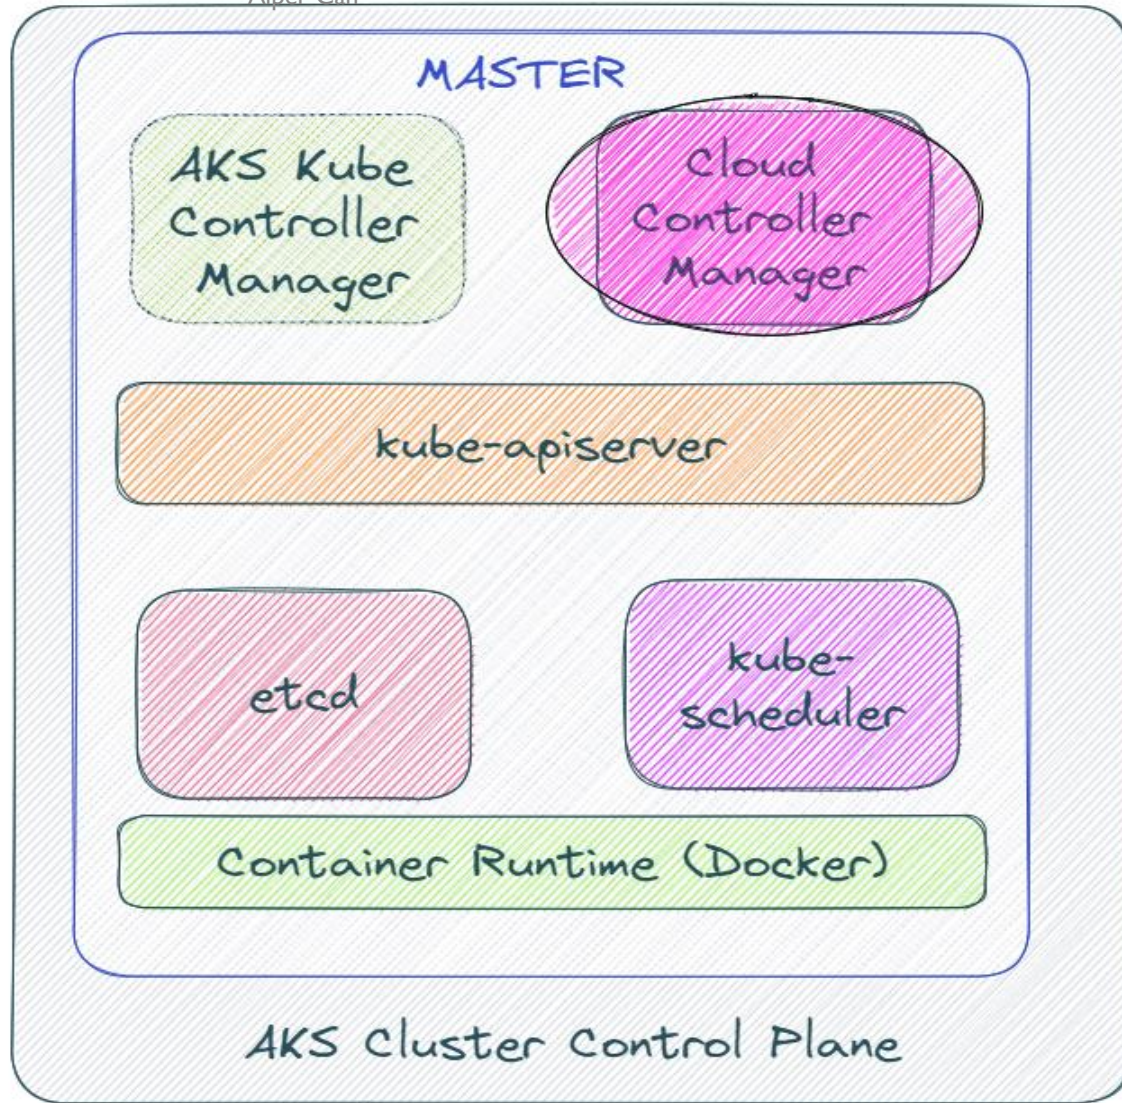
○ kube-scheduler

- Scheduler is responsible for distributing containers across multiple nodes.
- It watches for newly created Pods with no assigned node, and selects a node for them to run on.



○ kube-controller-manager

- Controllers are responsible for noticing and responding when nodes, containers or endpoints go down. They make decisions to bring up new containers in such cases.
- Node Controller: Responsible for noticing and responding when nodes go down.
- Replication Controller: Responsible for maintaining the correct number of pods for every replication controller object in the system.
- Endpoints Controller: Populates the Endpoints object (that is, joins Services & Pods)
- Service Account & Token Controller: Creates default accounts and API Access for new namespaces.



○ cloud-controller-manager

- A Kubernetes control plane component that embeds cloud-specific control logic.
- It only runs controllers that are specific to your cloud provider.
- On-Premise Kubernetes clusters will not have this component.
- Node controller: For checking the cloud provider to determine if a node has been deleted in the cloud after it stops responding.
- Route controller: For setting up routes in the underlying cloud infrastructure.
- Service controller: For creating, updating and deleting cloud provider load balancer.

Worker Node



○ Container Runtime

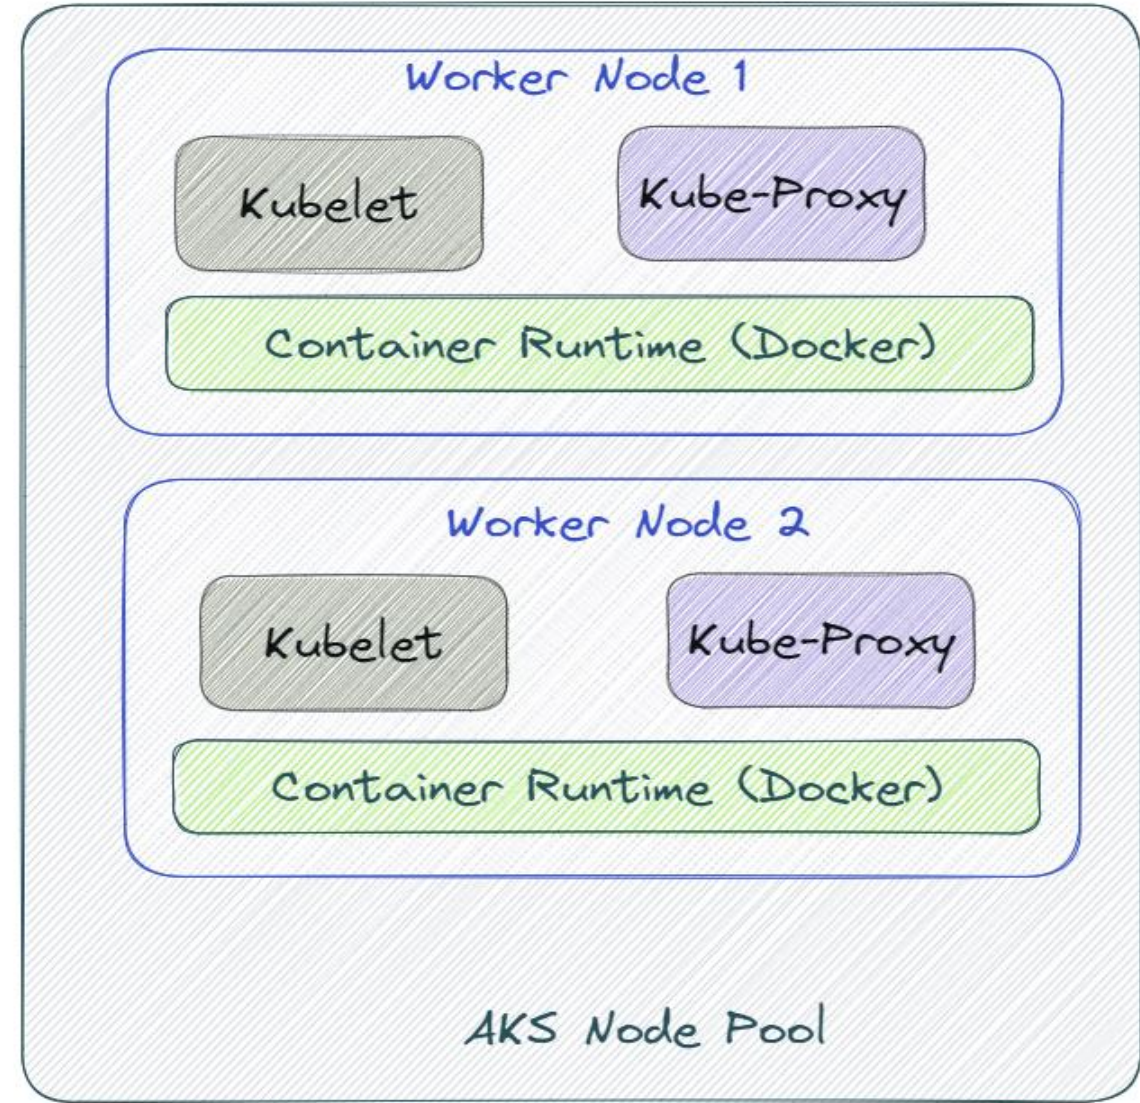
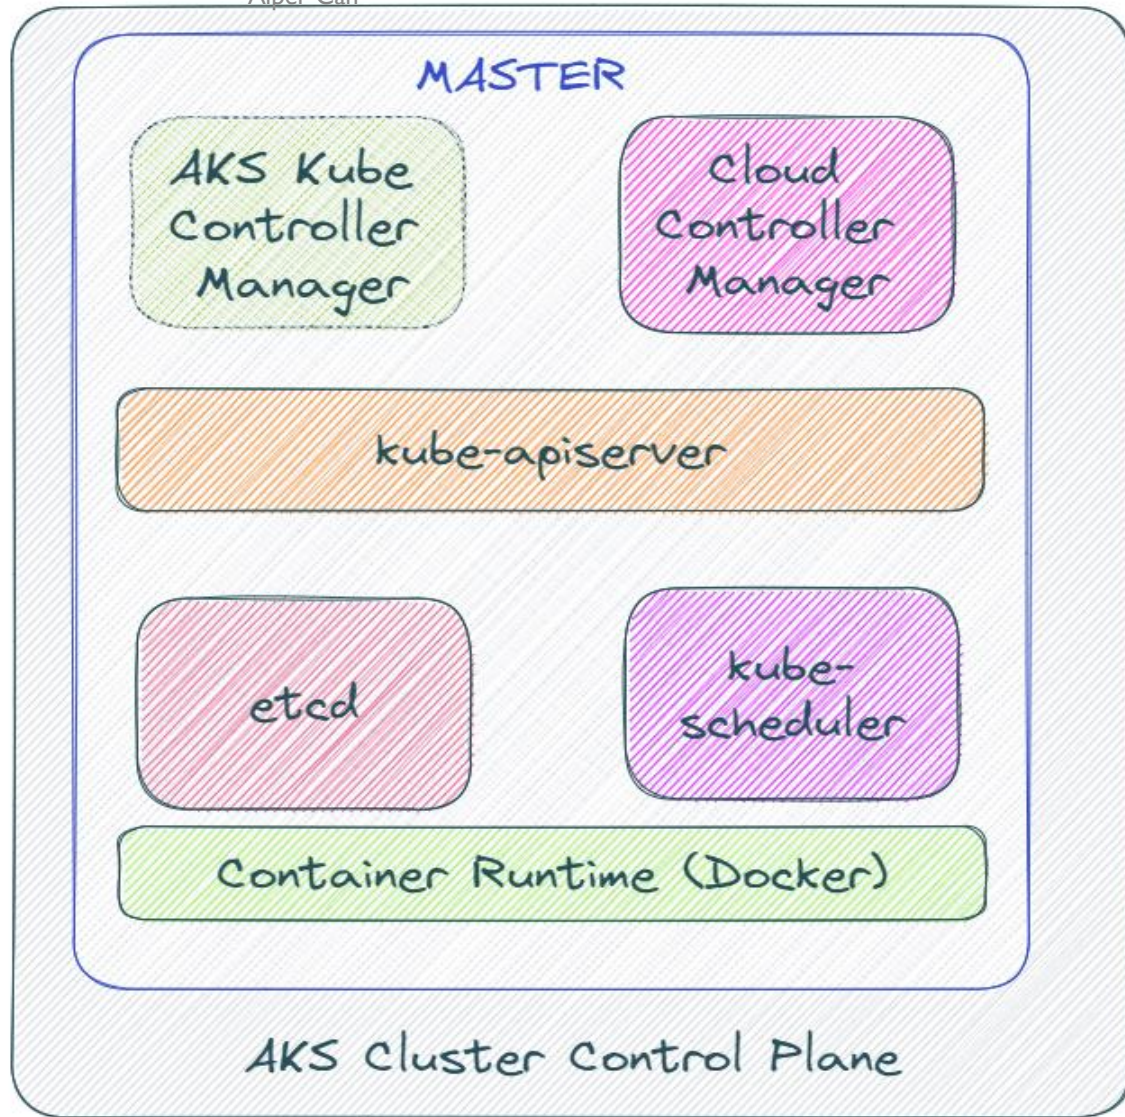
- Container Runtime is the underlying software where we run all these
- We are using Docker, but we have other runtime options like rkt, container-d etc.

○ Kubelet

- Kubelet is the agent that runs on every node in the cluster
- This agent is responsible for making sure that containers are running in a Pod on a node.

○ Kube-Proxy

- It is a network proxy that runs on each node in your cluster.
- It maintains network rules on nodes
- In short, these network rules allow network communication to your Pods from network sessions inside or outside of your cluster.



Kubernetes

Fundamentals

Pod, ReplicaSet, Deployment &
Service



KUBERNETES - FUNDAMENTALS

k8s Fundamentals

Pod

A POD is a single instance of an Application.
A POD is the smallest object, that you can create in Kubernetes.

ReplicaSet

A ReplicaSet will maintain a stable set of replica Pods running at any given time.
In short, it is often used to guarantee the availability of a specified number of identical Pods

Deployment

A Deployment runs multiple replicas of your application and automatically replaces any instances that fail or become unresponsive. Rollout & rollback changes to applications. Deployments are well-suited for stateless applications.

Service

A service is an abstraction for pods, providing a stable, so called virtual IP (VIP) address.
In simple terms, service sits Infront of a POD and acts as a load balancer.

KUBERNETES - IMPERATIVE & DECLARATIVE

Kubernetes Fundamentals

Imperative

kubectl

Pod

ReplicaSet

Deployment

Service

Declarative

YAML & kubectl

Pod

ReplicaSet

Deployment

Service

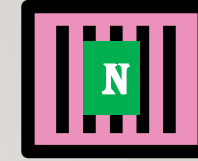
Kubernetes

POD

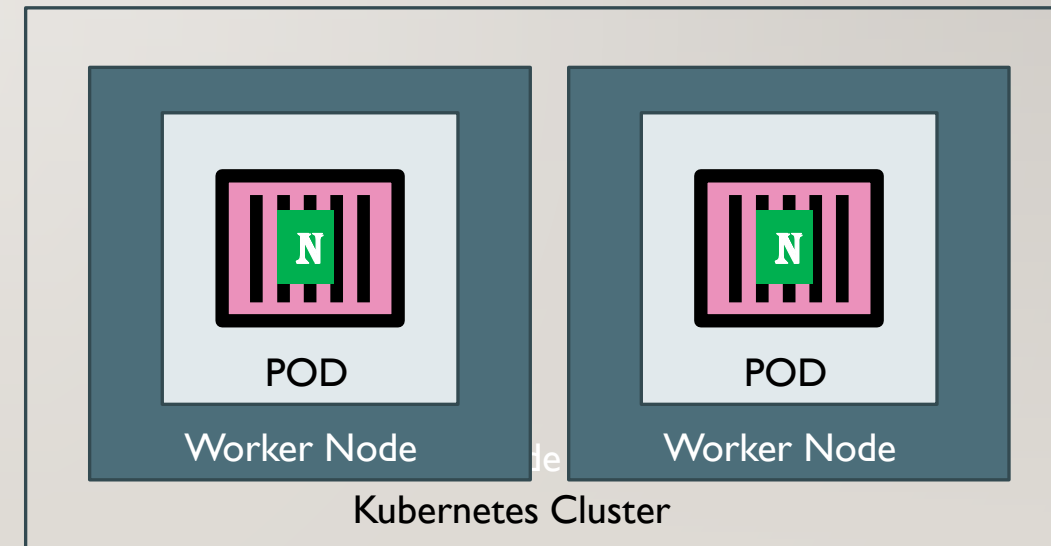


KUBERNETES - POD

- With Kubernetes our core goal will be to **deploy our applications** in the form of **containers** on **worker nodes** in a k8s cluster.
- Kubernetes **does not** deploy containers directly on the worker nodes.
- Container is **encapsulated** in to a Kubernetes Object named **POD**.
- A POD is a **single instance** of an application.
- A POD is the **smallest object** that we can create in Kubernetes.

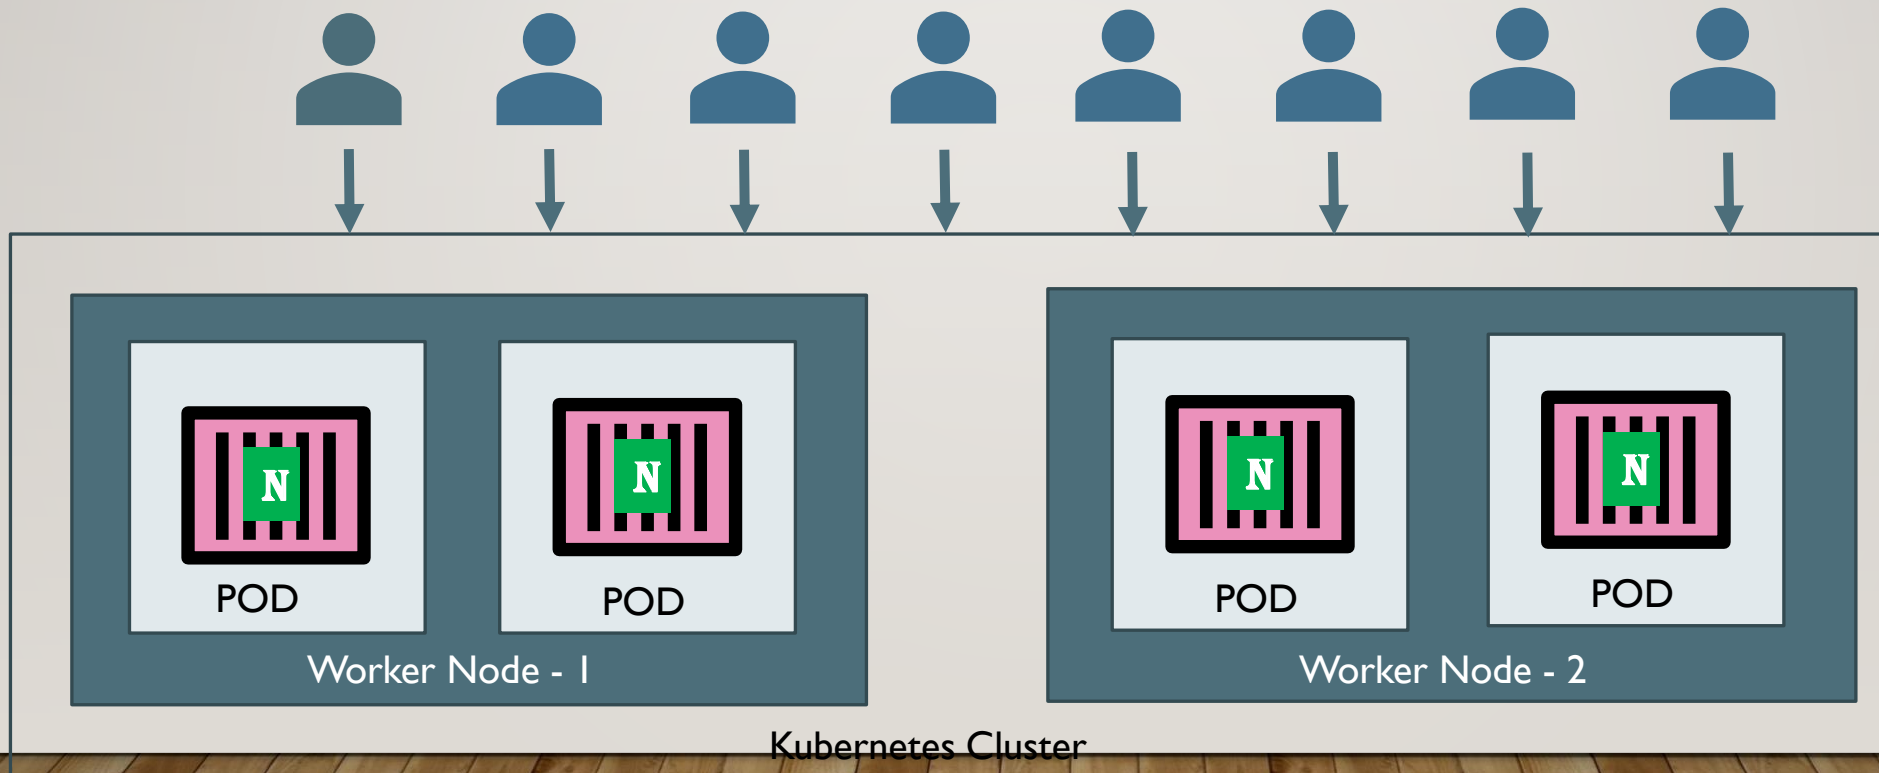


Nginx Container
Image



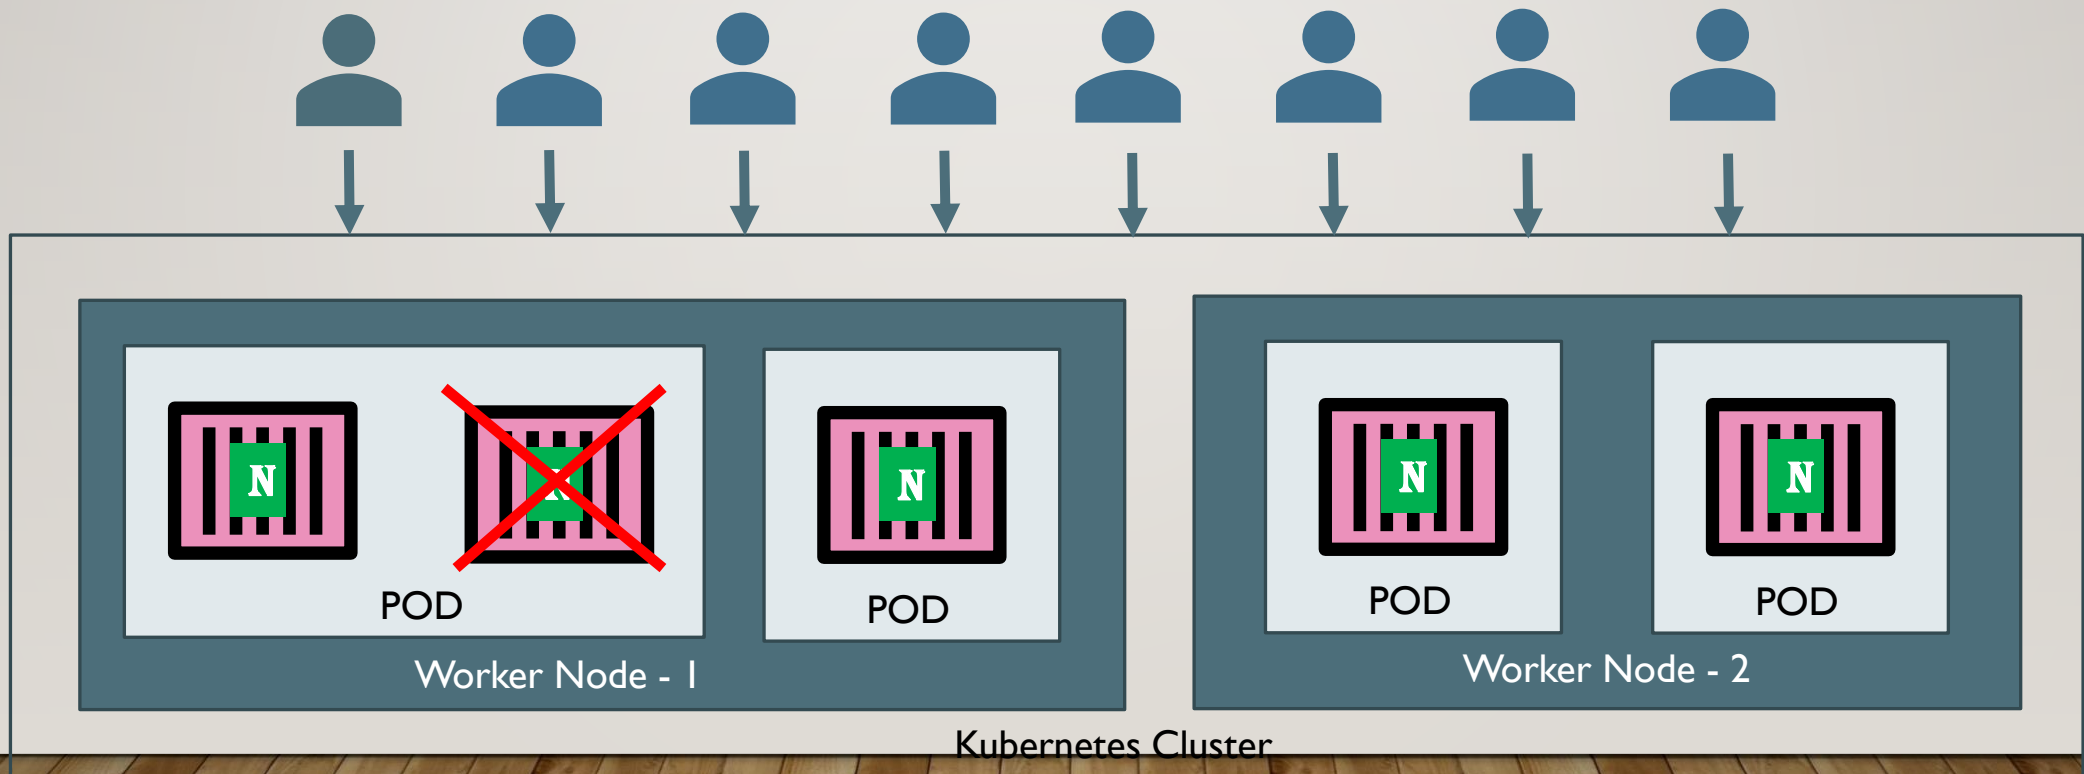
KUBERNETES - POD

- PODs generally have **one to one** relationship with containers.
- To scale up we **create** new POD and to scale down we **delete** the POD.



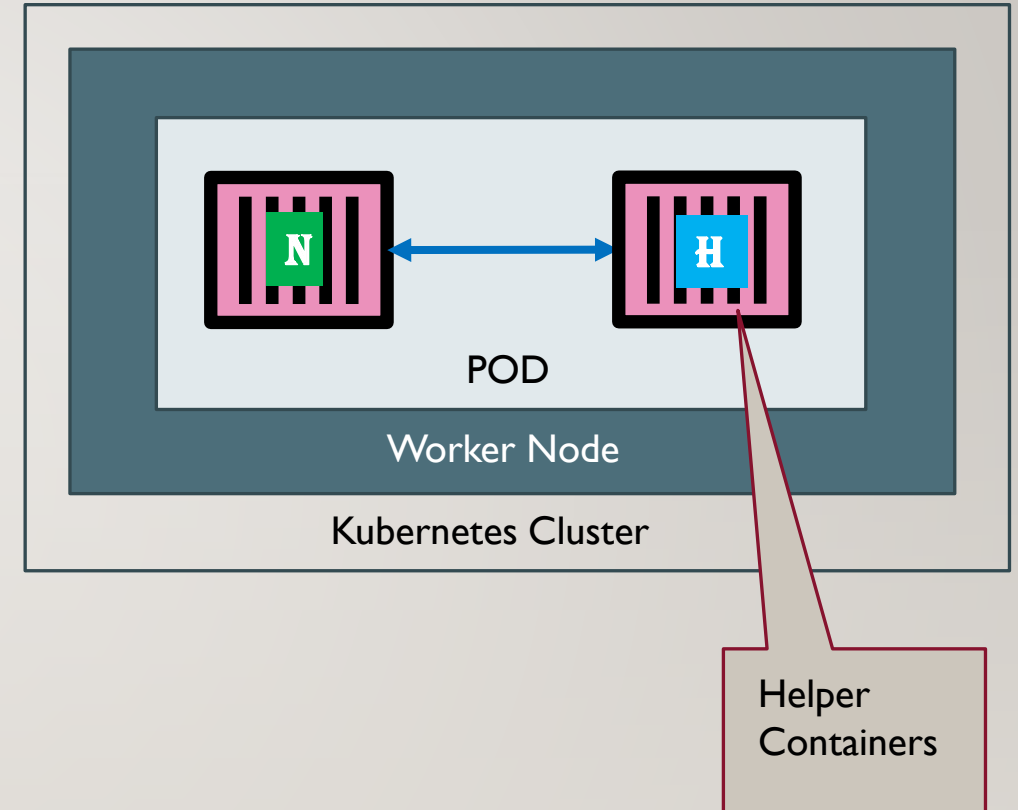
KUBERNETES – PODS

- We **cannot have** multiple containers of **same kind** in a single POD.
- Example: Two NGINX containers in single POD serving same purpose is **not recommended**.



KUBERNETES – MULTI-CONTAINER PODS

- We can have multiple containers in a single POD, provided they are not of same kind.
- Helper Containers (Side-car)
 - Data Pullers: Pull data required by Main Container
 - Data pushers: Push data by collecting from main container (logs)
 - Proxies: Writes static data to html files using Helper container and Reads using Main Container.
- Communication
 - The two containers can easily communicate with each other easily as they share same **network space**.
 - They can also easily share **same storage space**.
- Multi-Container Pods is a **rare use-case** and we will try to focus on core fundamentals.

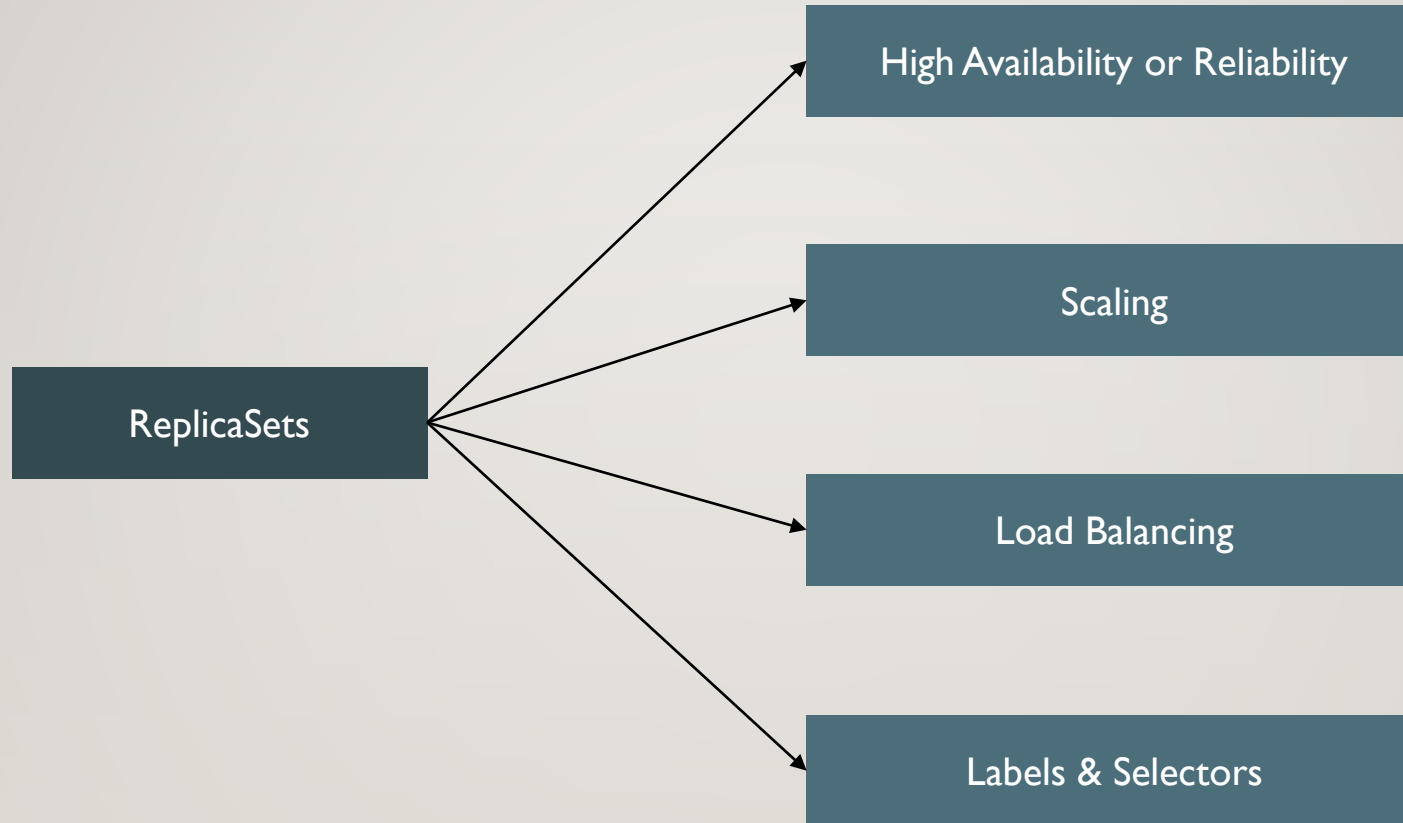


Kubernetes

ReplicaSets



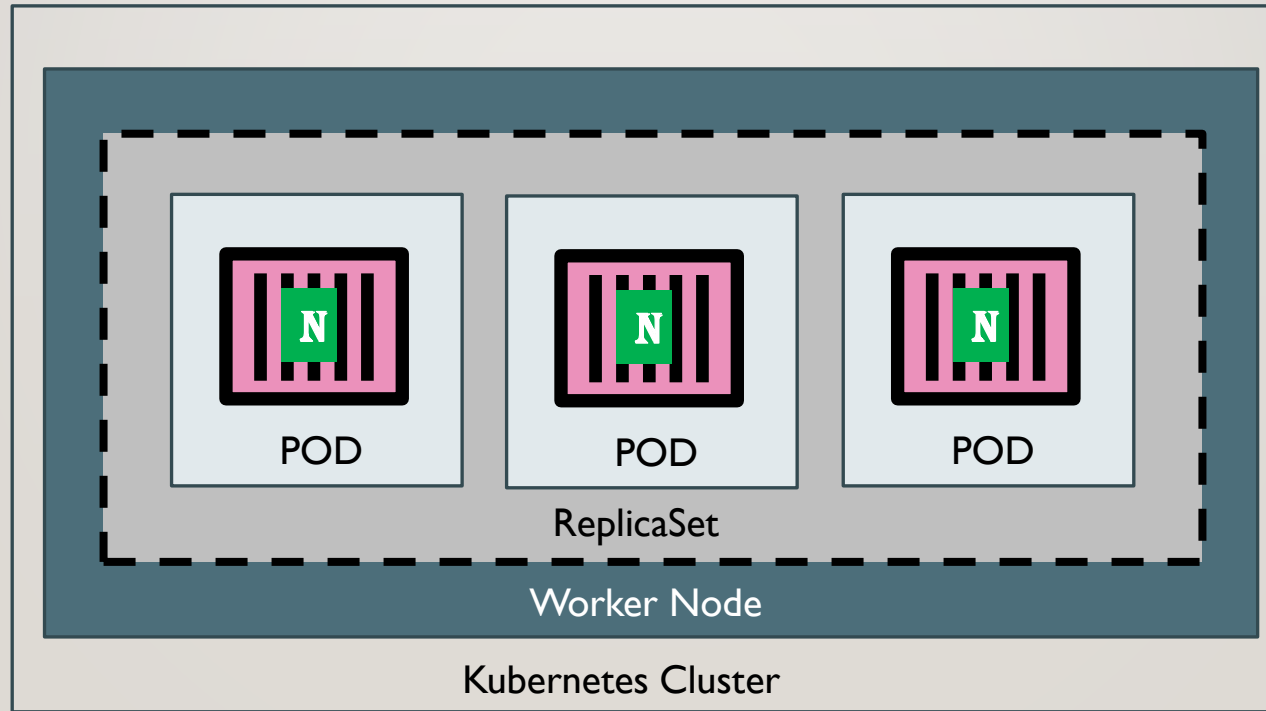
KUBERNETES - REPLICASETS



KUBERNETES – REPLICASET

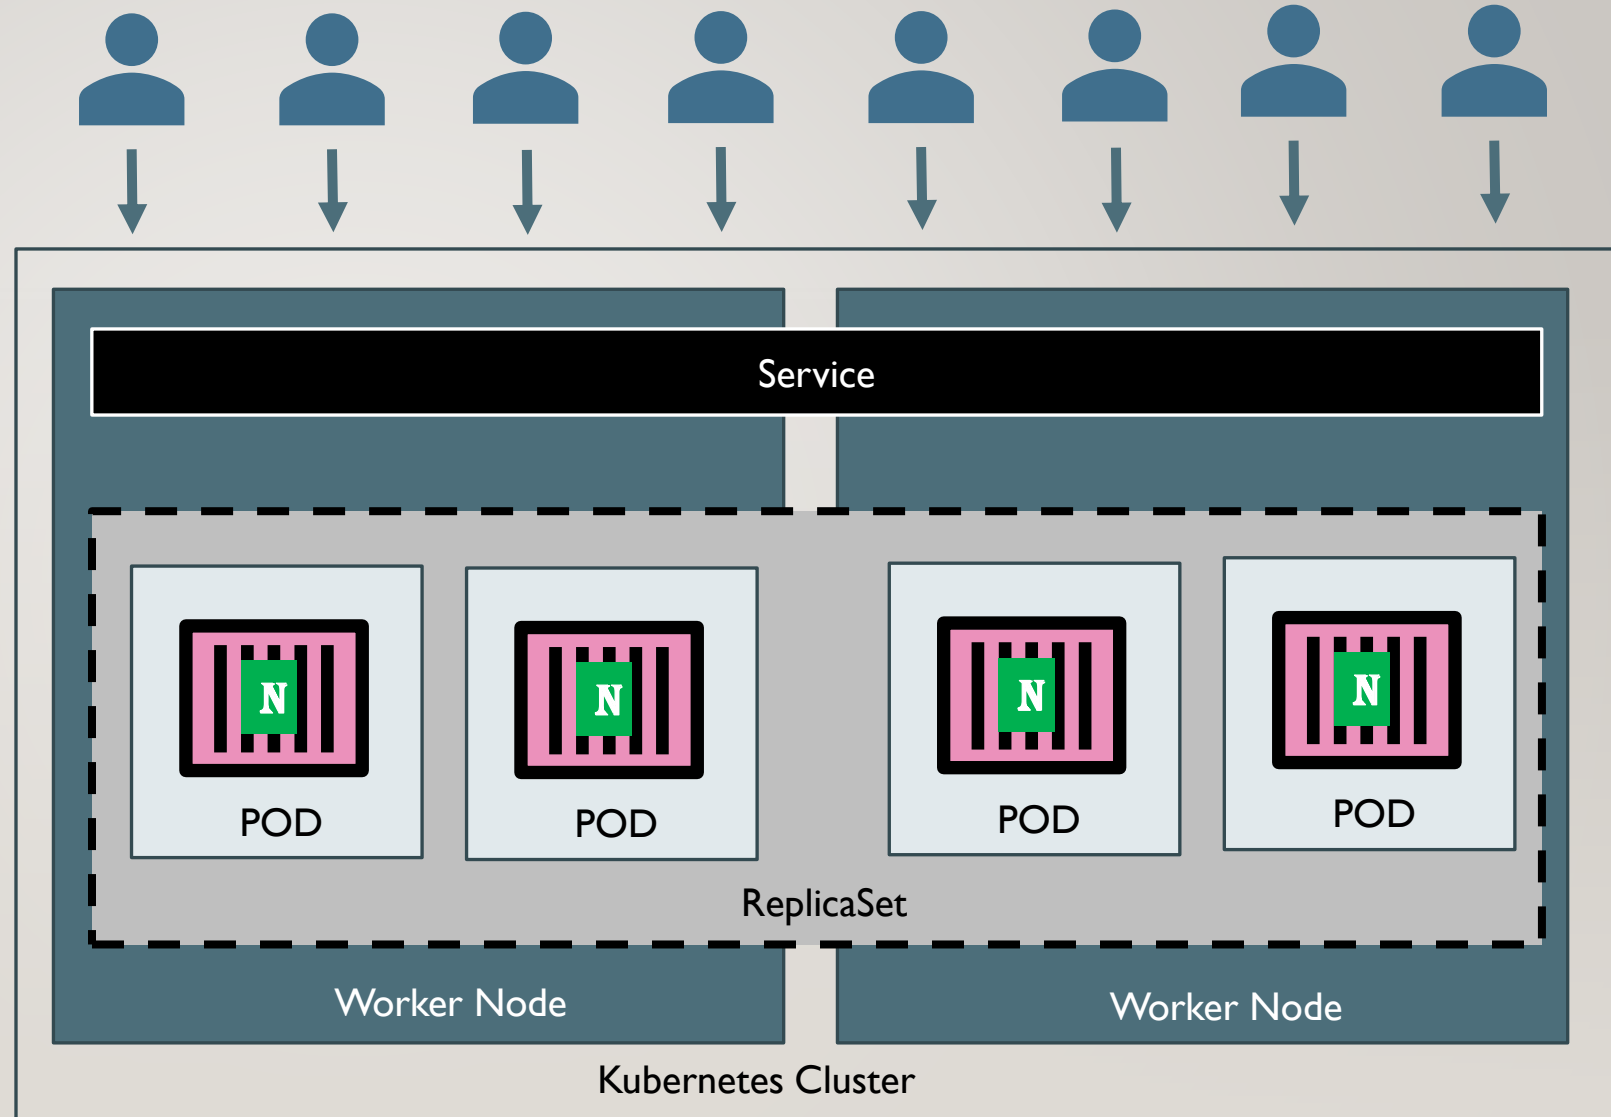
- A ReplicaSet's purpose is to maintain a **stable set of replica Pods** running at any given time.
- If our **application crashes (any pod dies)**, replicaset will **recreate** the pod immediately to ensure the configured number of pods running at any given time.

Reliability
Or
High Availability



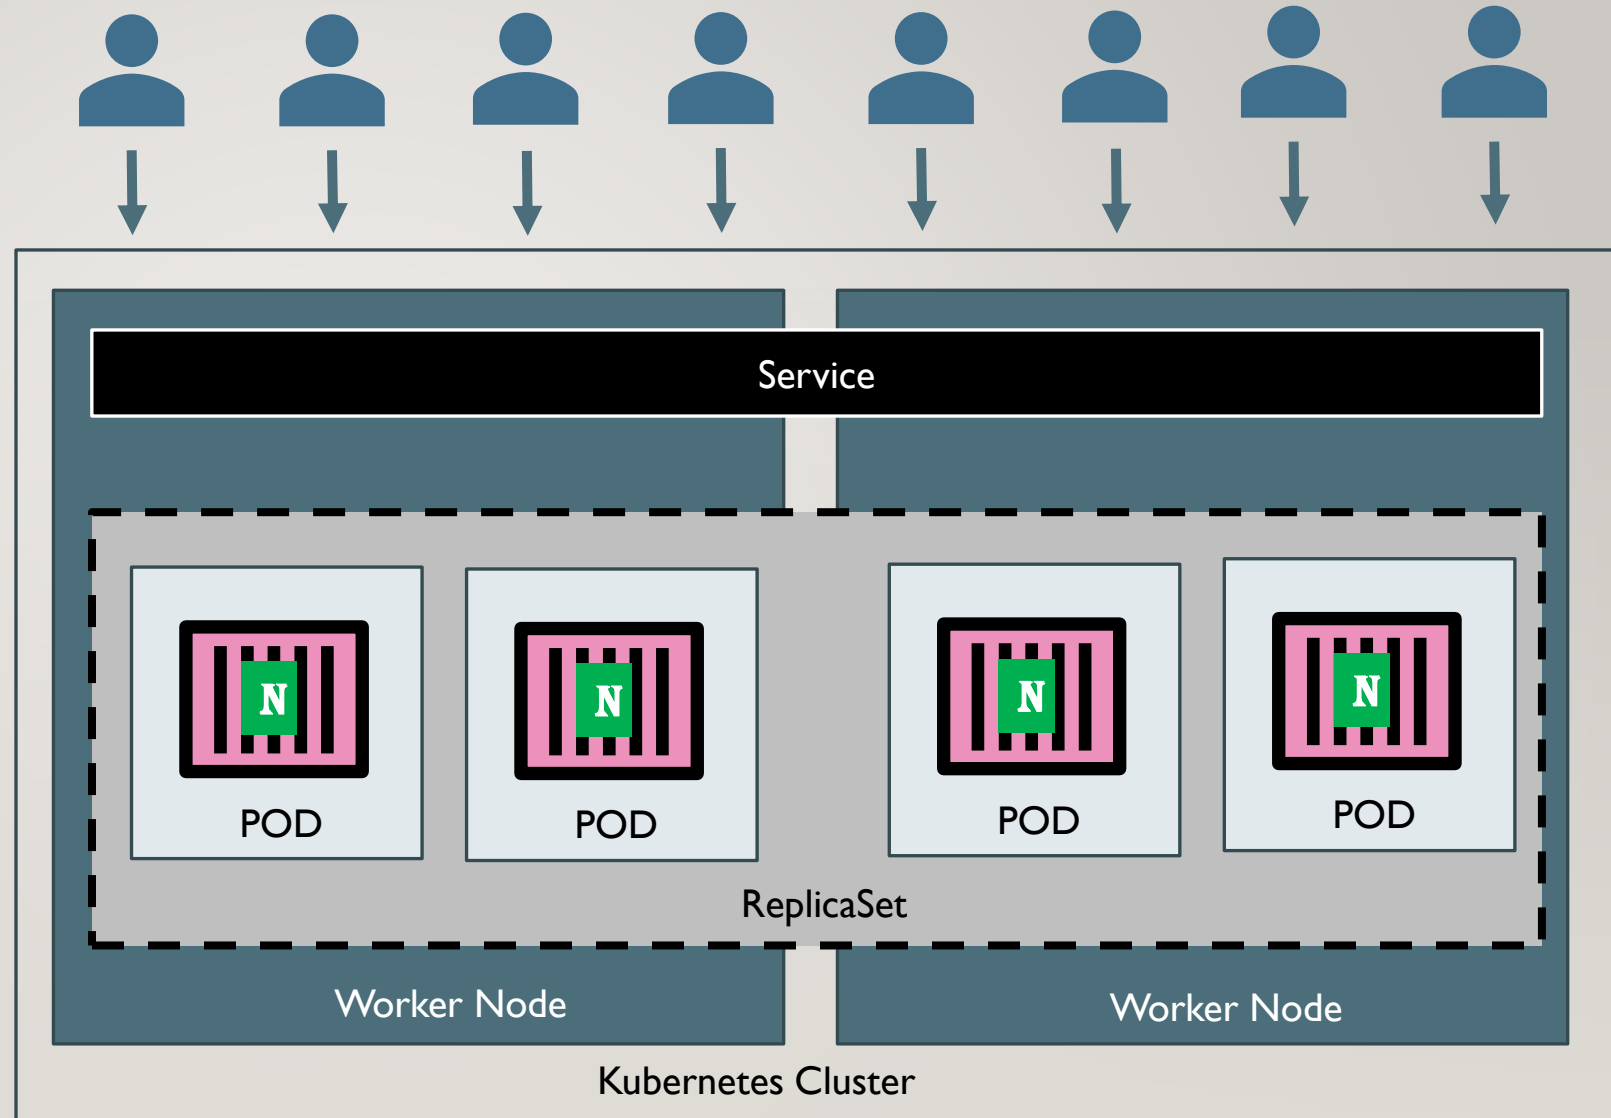
KUBERNETES – REPLICASET

- Load Balancing
- To avoid overloading of traffic to single pod we can use **load balancing**.
- Kubernetes provides pod load balancing **out of the box** using **Services** for the pods which are part of a ReplicaSet
- **Labels & Selectors** are the **key items** which **ties** all 3 together (Pod, ReplicaSet & Service), we will know in detail when we are writing YAML manifests for these objects



KUBERNETES – REPLICASET

- Scaling
- When load become too much for the number of existing pods, Kubernetes enables us to easily **scale** up our application, adding additional pods as needed.
- This is going to be **seamless and super quick**.



Kubernetes

ReplicaSets

Demo



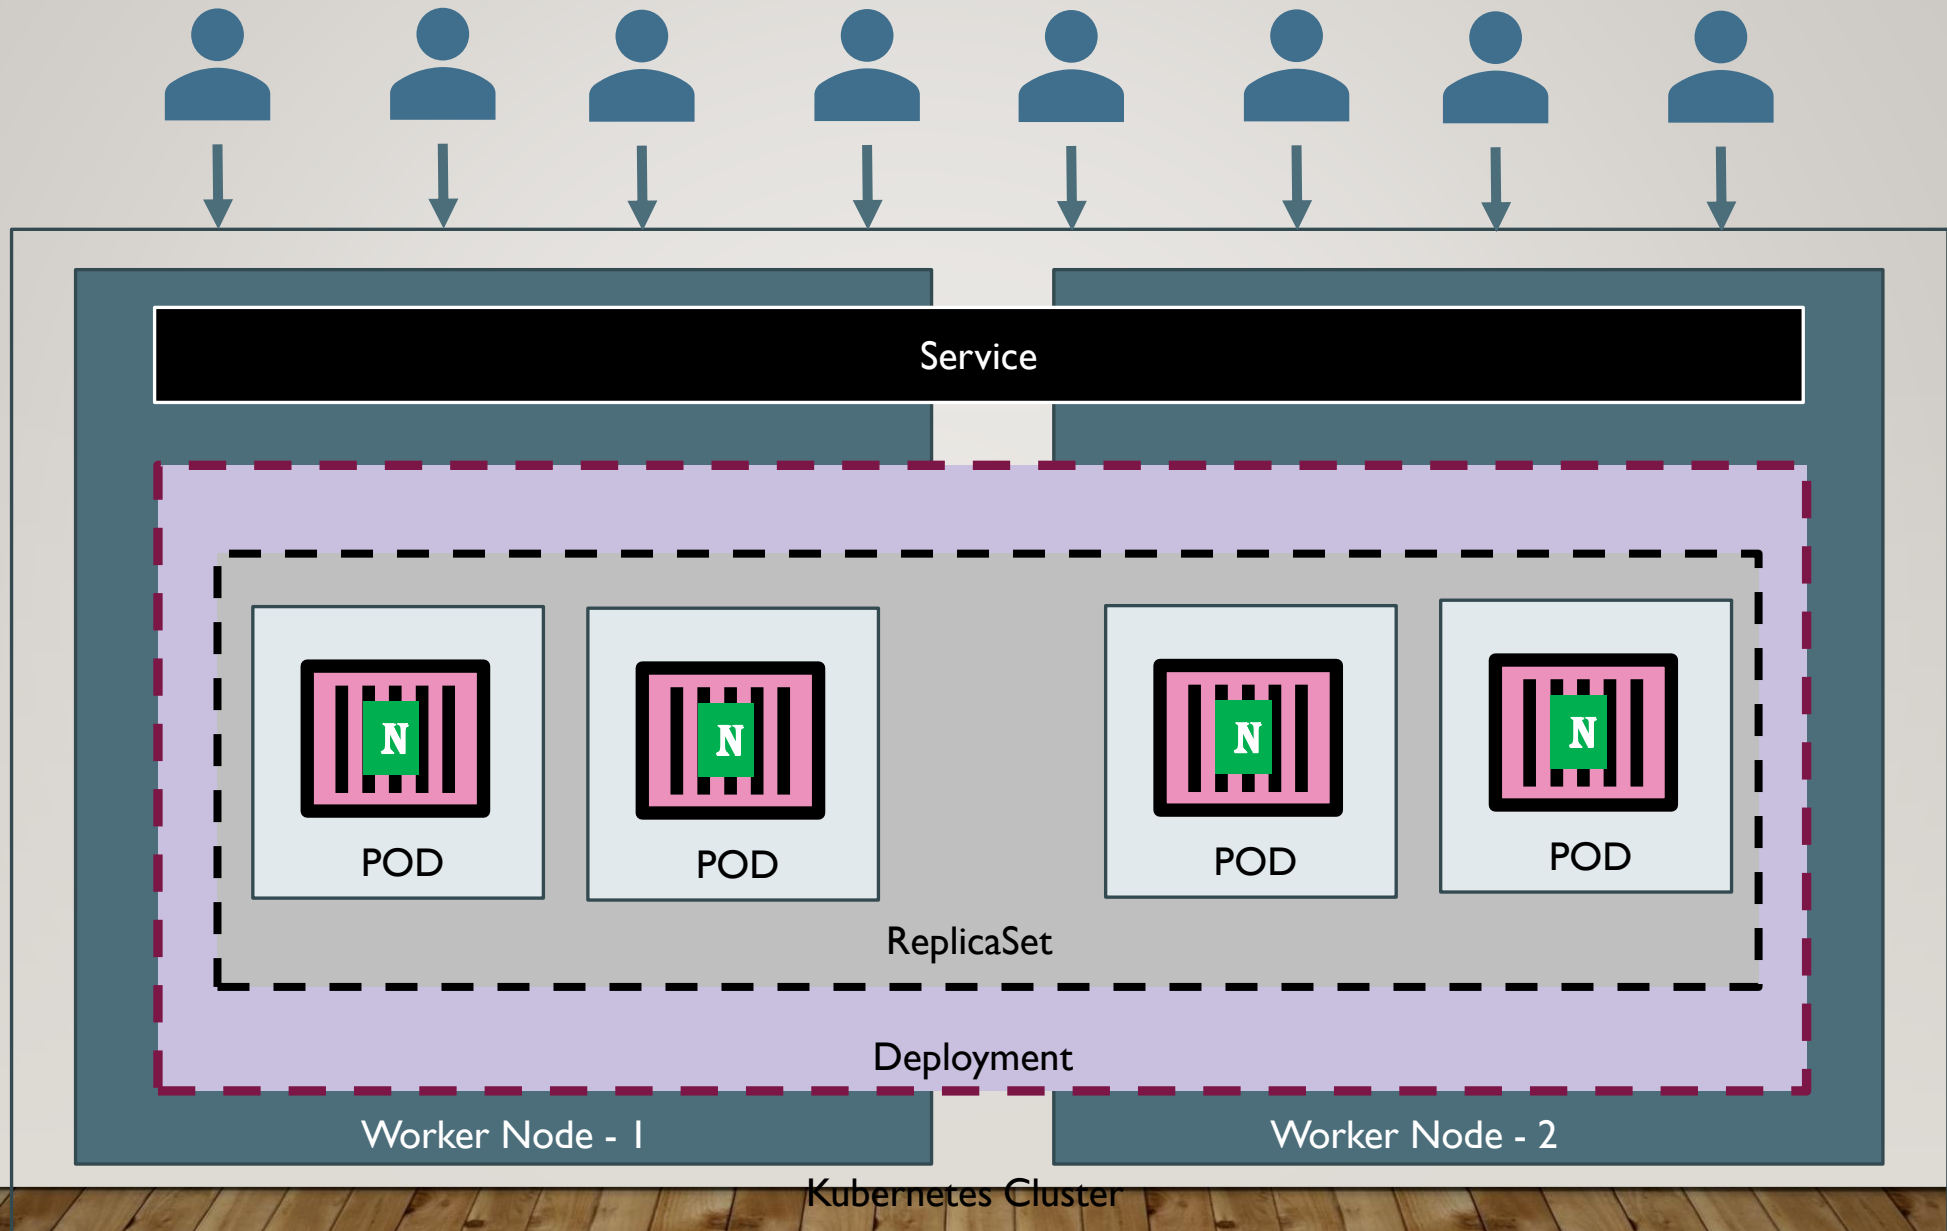
Kubernetes Deployments



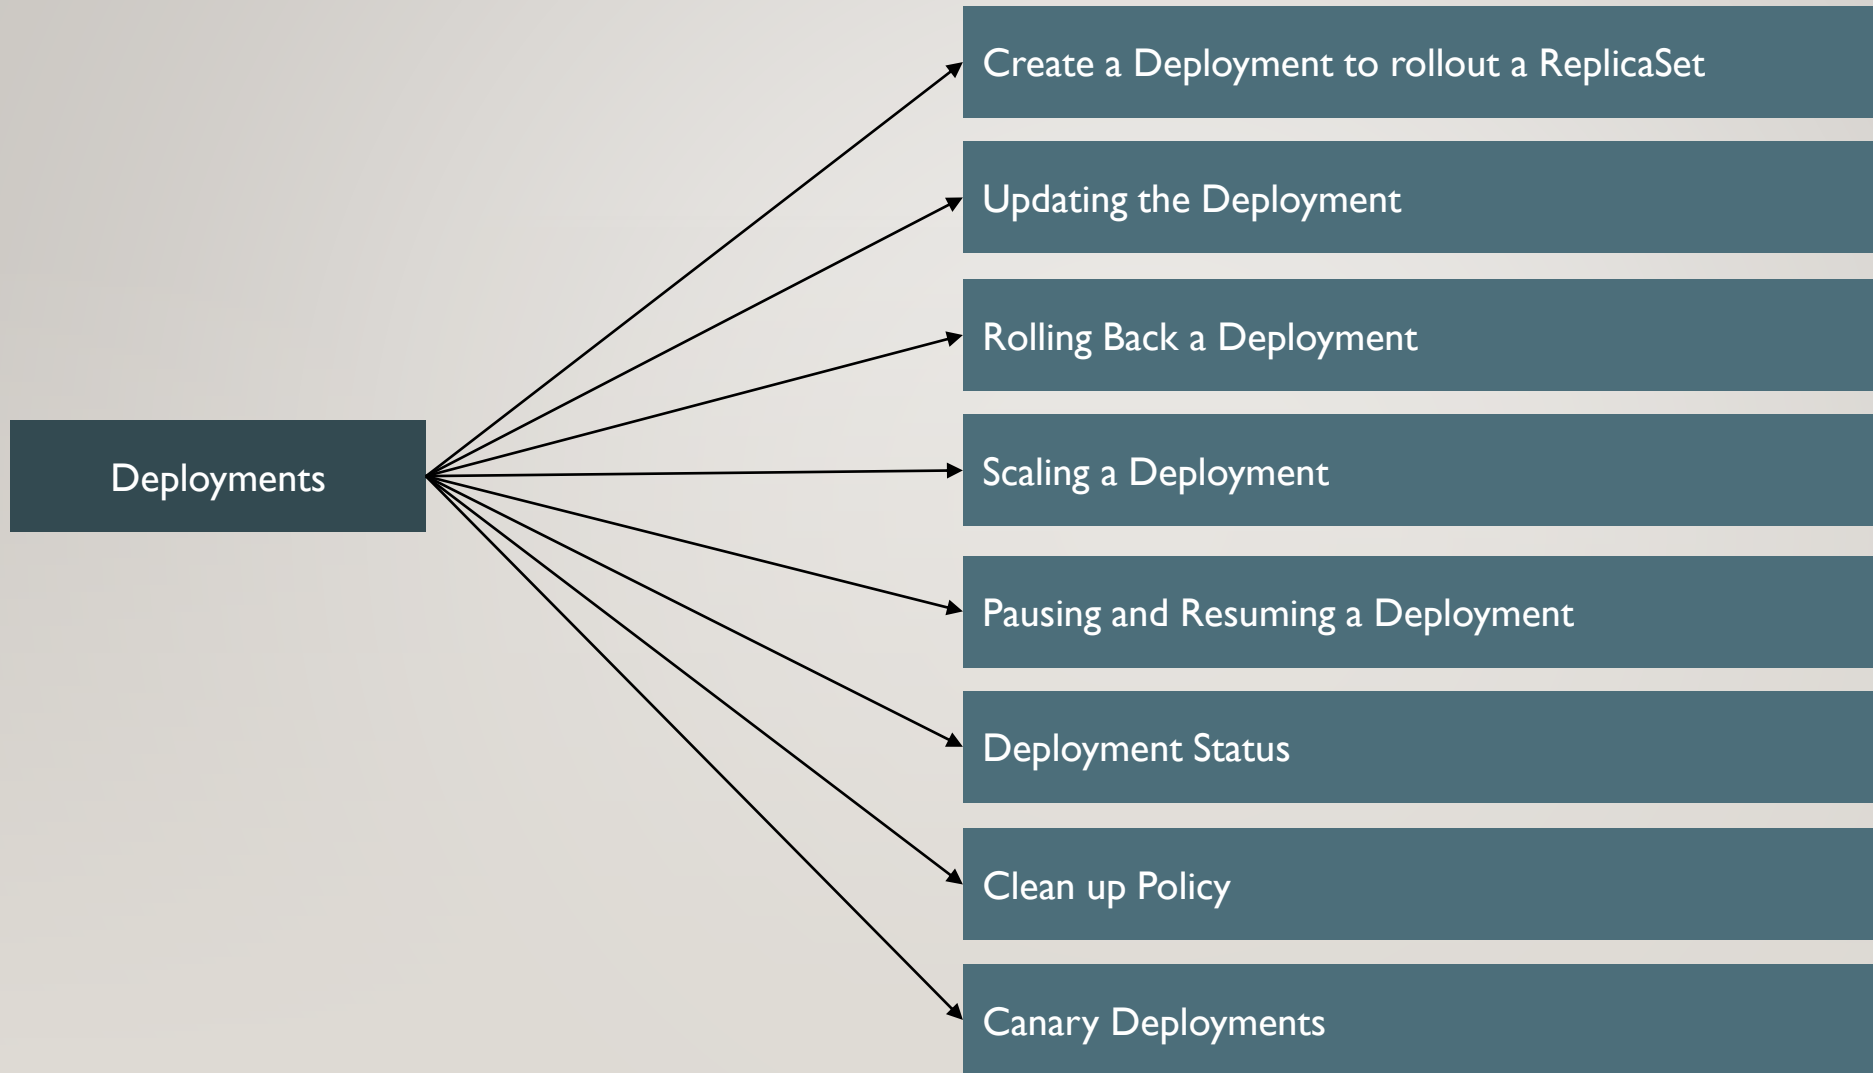
Kubernetes Deployments



KUBERNETES – DEPLOYMENTS



KUBERNETES - DEPLOYMENT



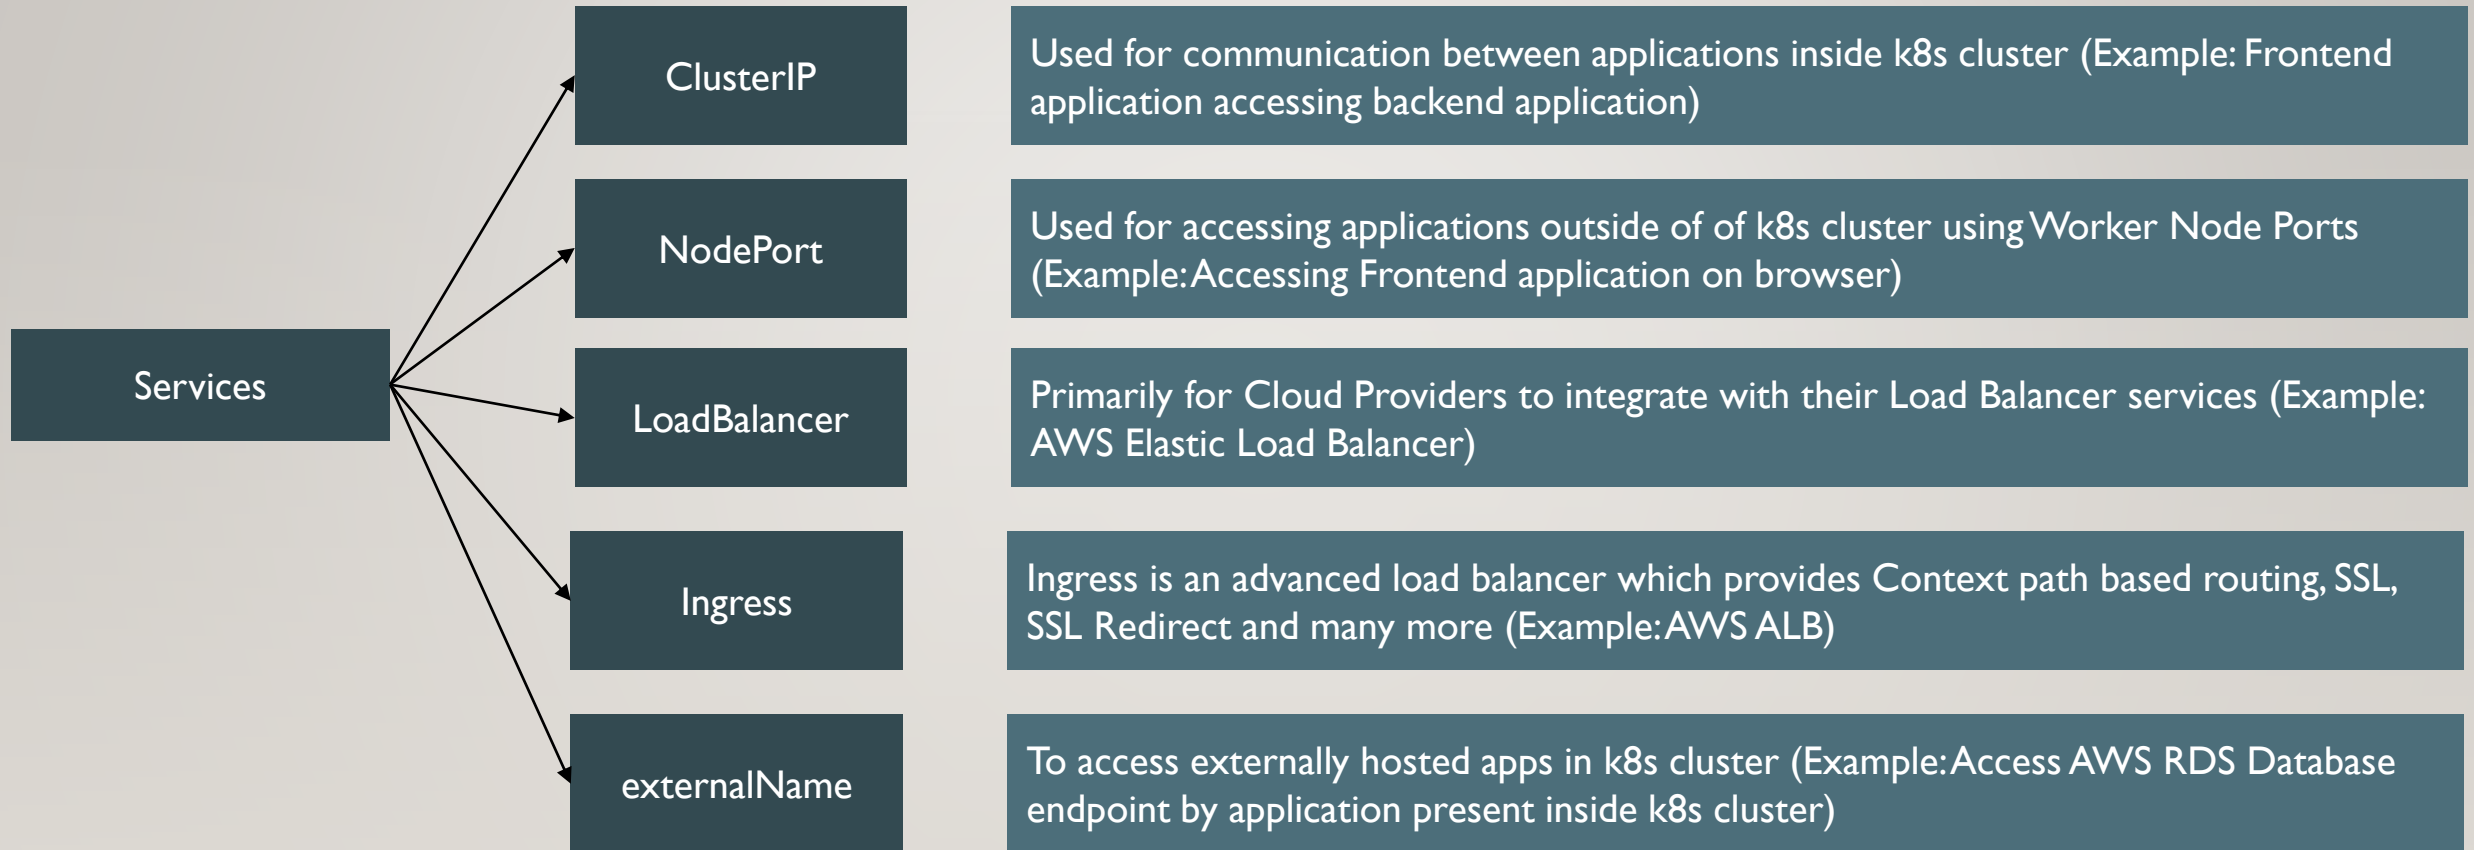
Kubernetes Deployments Demo



Kubernetes Services



KUBERNETES - SERVICES

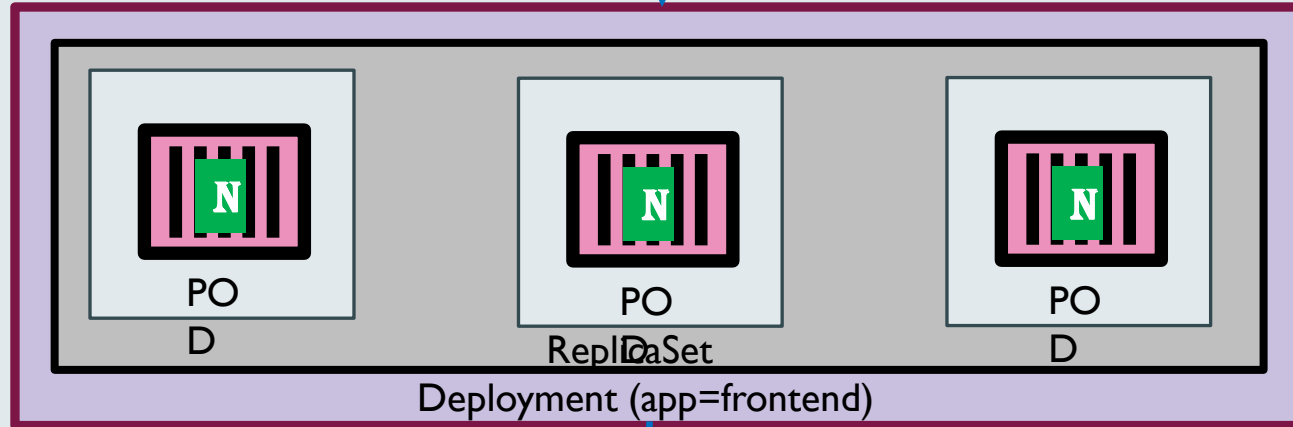




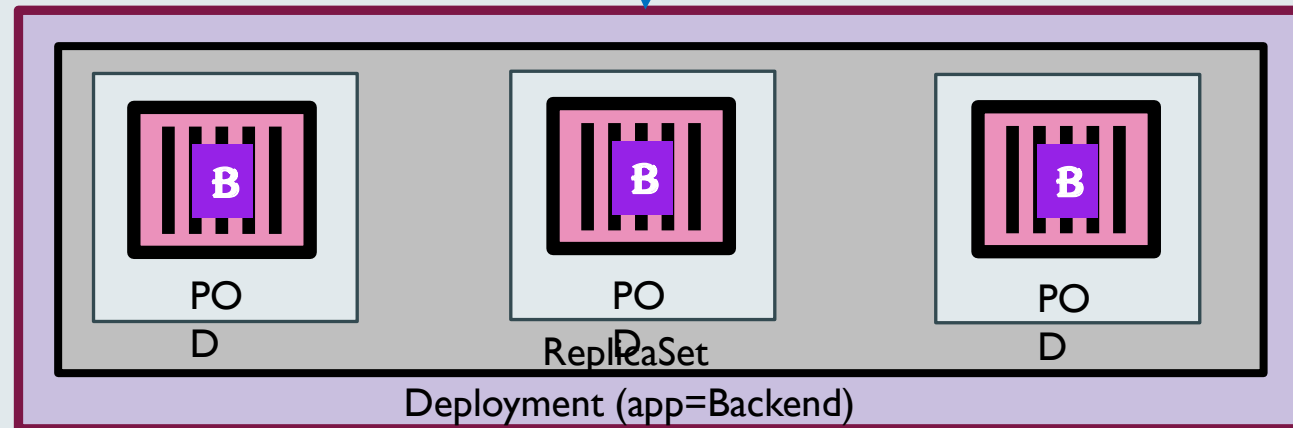
Users

Kubernetes Cluster

Frontend App – NodePort or LoadBalancer or Ingress Service



Backend App - ClusterIP Service



DB – ExternalName Service



SERVICES

Kubernetes

Services

Demo



Kubernetes

YAML Basics



YAML BASICS

- YAML is used to **store information** about different things
- We can use YAML to **define key, Value pairs** like variables, lists and objects
- YAML is very similar to **JSON** (Javascript Object Notation)
- YAML primarily focuses on **readability** and **user friendliness**
- YAML is designed to be **clean and easy to read**
- We can define YAML files with two different extensions
 - **abc.yml**
 - **abc.yaml**

YAML BASICS

- YAML Comments
- YAML Key Value Pairs
- YAML Dictionary or Map
- YAML Array / Lists
- YAML Spaces
- YAML Document Separator





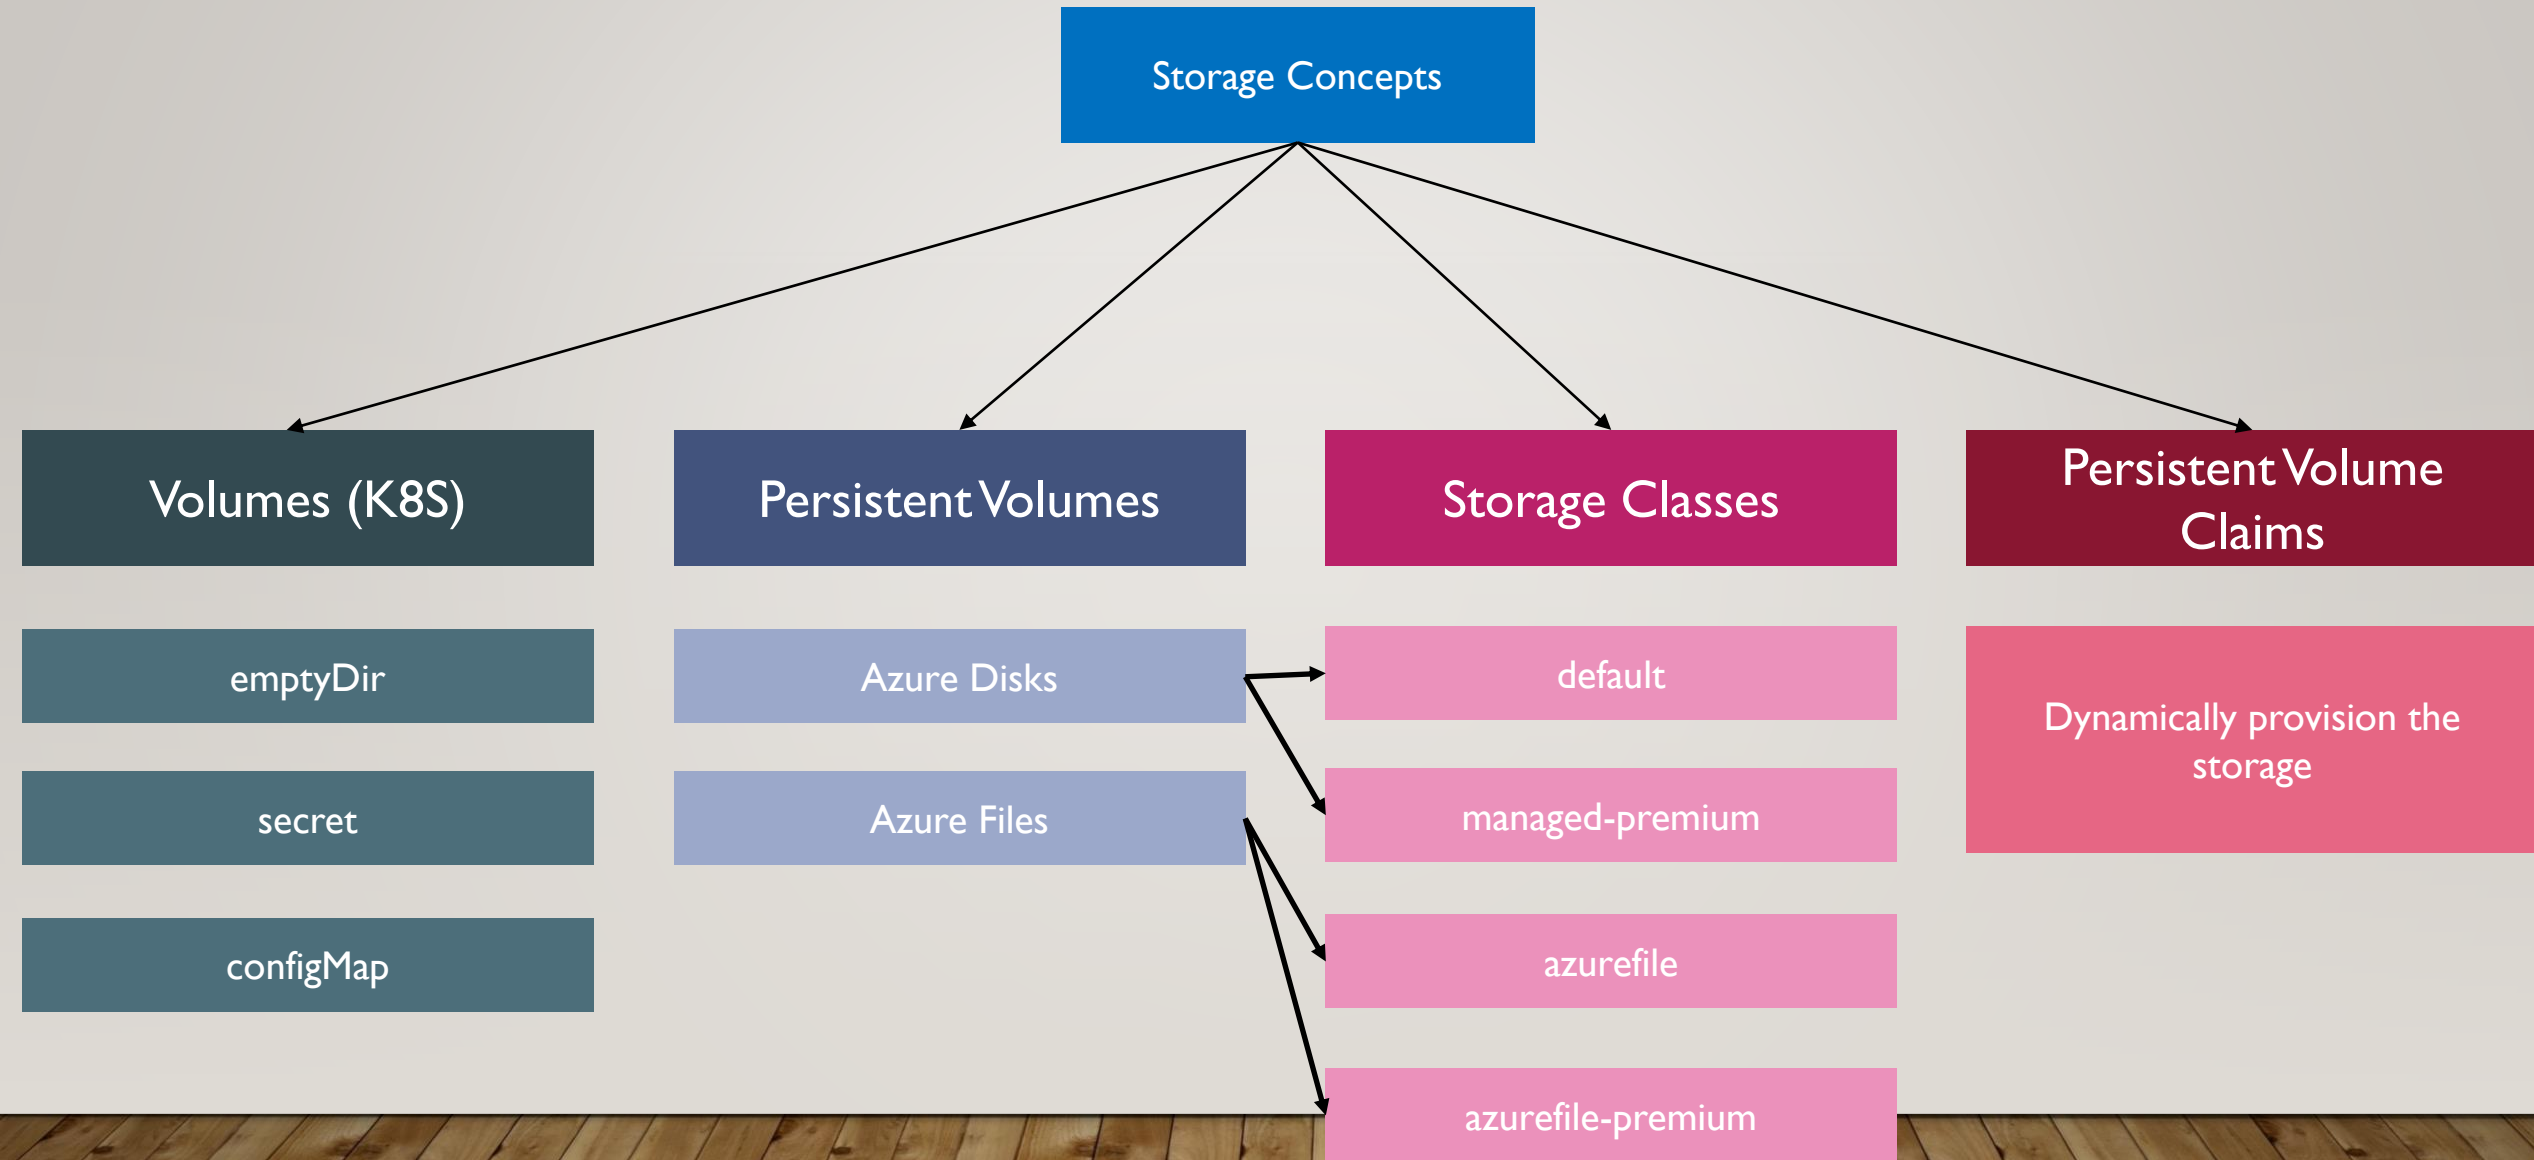
Azure AKS Storage

with

Azure Disks & Files

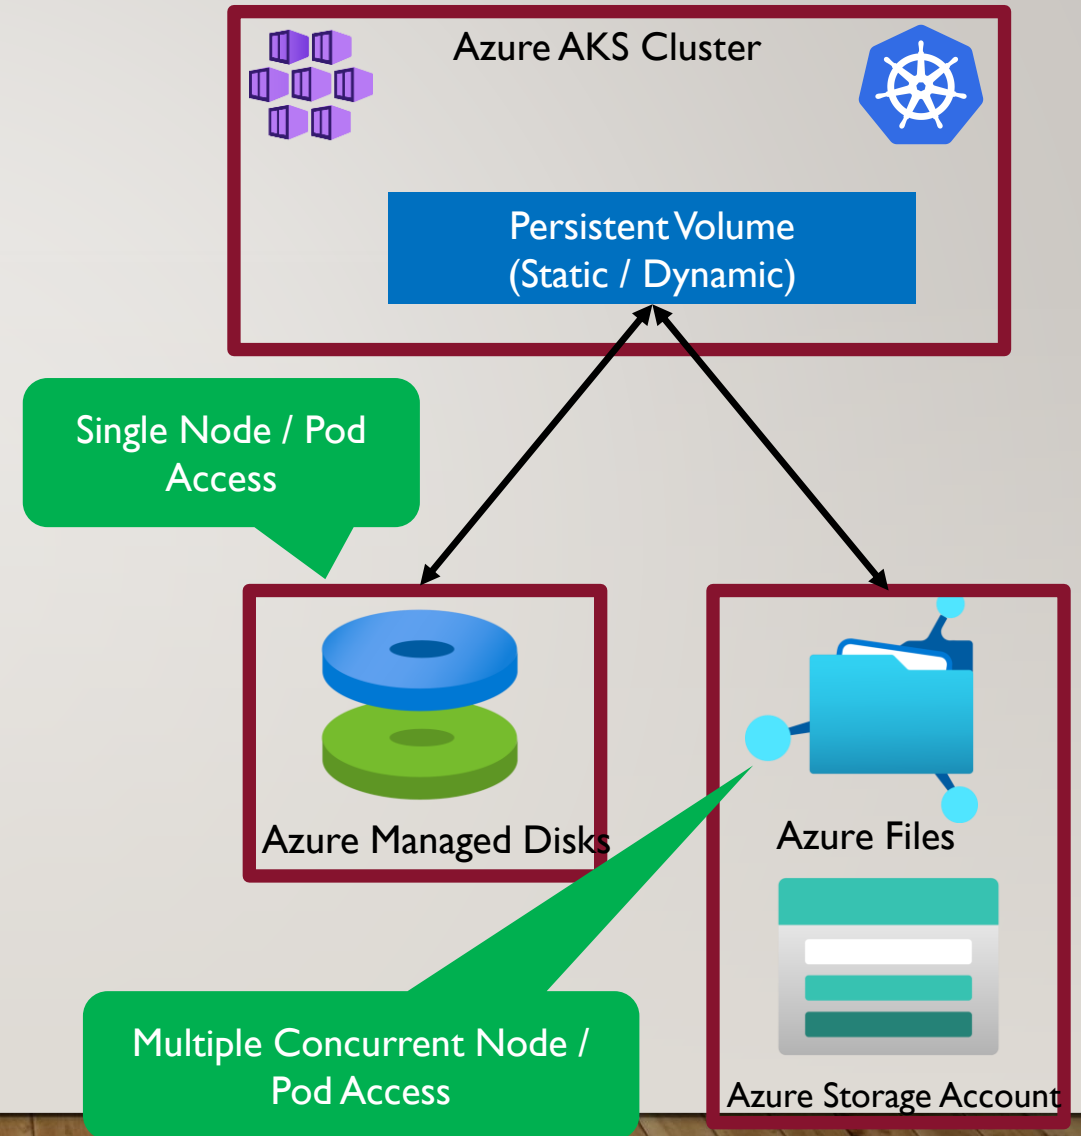


AKS STORAGE CONCEPTS



WHY DO WE NEED PERSISTENT STORAGE IN KUBERNETES?

- Volumes that are defined and created as part of the **pod lifecycle** only exist until the pod is **deleted**.
- Pods often expect their storage to remain if a pod is **rescheduled** on a different host during a maintenance event, especially in StatefulSets.
- A **persistent volume (PV)** is a storage resource created and managed by the Kubernetes API that can exist beyond the lifetime of an individual pod.
- For AKS, **Azure Disks** or **Azure Files** are used to provide the PersistentVolume.



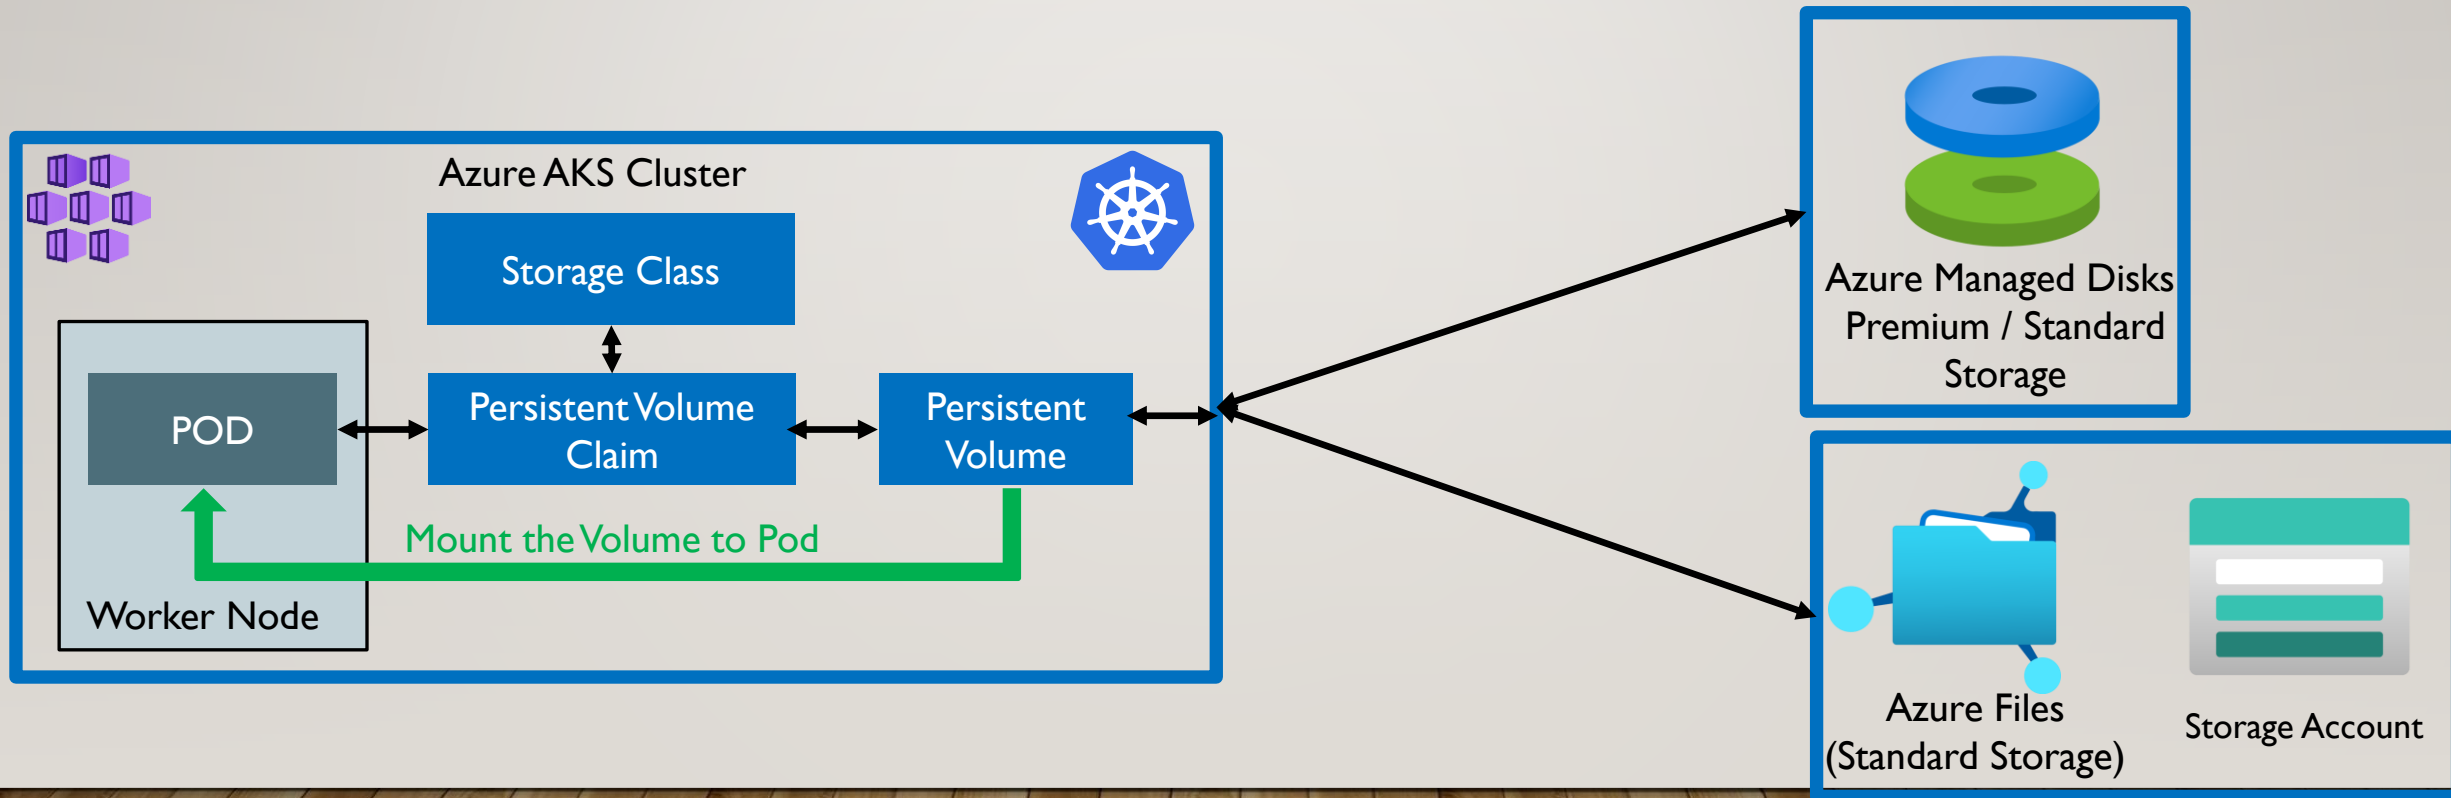
PERSISTENT VOLUME CLAIMS – HOW IT WORKS?

create a pod

PVC will request storage based on specifics from Storage class (Disk or File)

Kube API Server will dynamically provision the Persistent Volume PV requested

Once the PV is available it will get mounted to Pod for usage.





Azure AKS Storage

with

Azure Disks



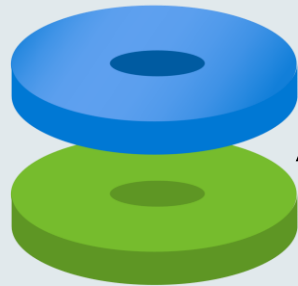
AZURE DISKS - INTRODUCTION

- Azure Disk Storage offers **high-performance, highly durable** block storage for our mission- and business-critical workloads
- We can mount these **volumes as devices** on our Virtual Machines & Container instances.
- **Cost-effective storage**
 - Built-in bursting capabilities to handle unexpected traffic and process batch jobs cost-effectively
- **Unmatched resiliency**
 - 0 percent annual failure rate for consistent enterprise-grade durability
- **Seamless scalability**
 - Dynamic scaling of disk performance on Ultra Disk Storage without disruption
- **Built-in security**
 - Automatic encryption to help protect your data using Microsoft-managed keys or your own



AKS STORAGE AZURE DISKS

Azure AKS Cluster



Azure Disks



Storage Class



Persistent Volume Claim



Persistent Volumes



Deployment

MySQL – ClusterIP Service



PO

ReplicaSet

Deployment (mysql)



Config Map



Environment Variables



Volumes



Volume Mounts

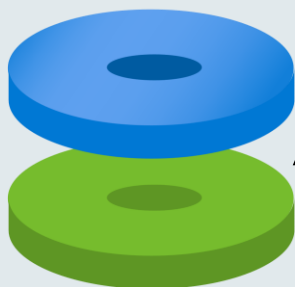


ClusterIP Service



Users

<http://<External-IP-from-get-service-output>>



Azure Disks



Storage Class



Persistent Volume Claim

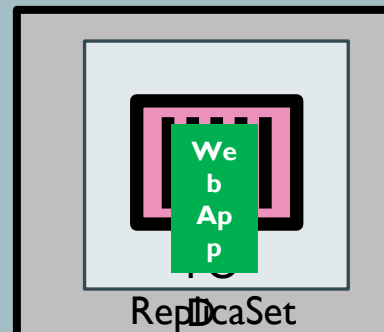


Persistent Volumes



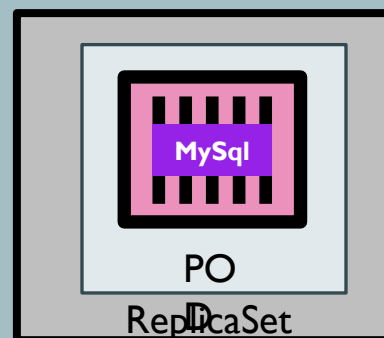
Deployment

UserMgmt – LoadBalancer Service



Deployment (UserMgmt)

MySQL – ClusterIP Service



Deployment (mysql)

Azure AKS Cluster



Load Balancer Service



Deployment



Environment Variables



Init Containers



Config Map



Environment Variables



Volumes



Volume Mounts



ClusterIP Service

AKS
Storage

Azure Disks



Azure AKS Storage

with

**Azure MySQL
Database**





Users

<http://<External-IP-from-get-service-output>>

Storage Azure MySQL Database

Drawbacks of Azure Disks

Complex setup to achieve HA, Statefulsets

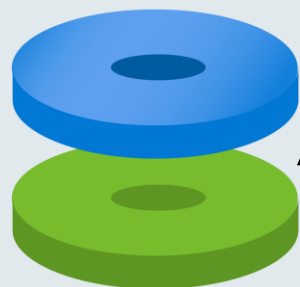
Complex Master-Master MySQL setup

Complex Master-Slave MySQL setup

No Automatic Backup & Recovery

No Auto-Upgrade MySQL

Logging, Monitoring needs custom scripts



Azure Disks



Storage Class



Persistent Volume Claim

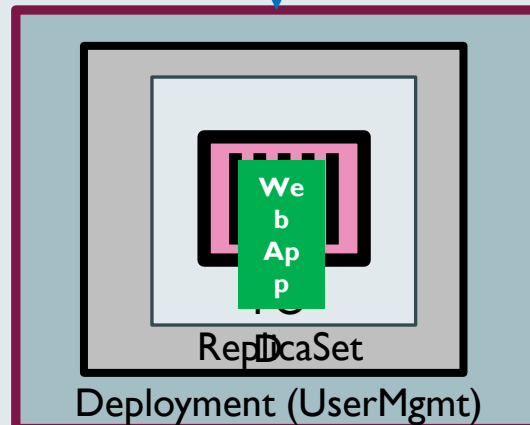


Persistent Volumes

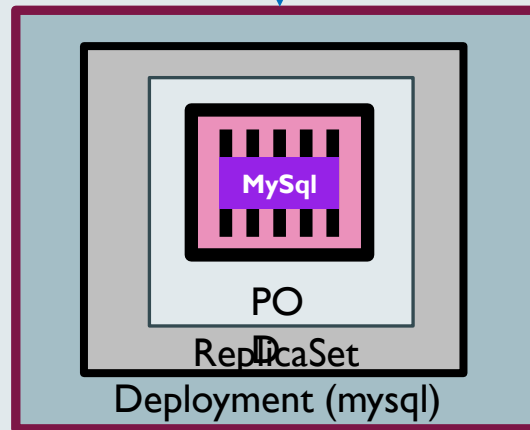


Deployment

UserMgmt – LoadBalancer Service



MySQL – ClusterIP Service



Azure AKS Cluster



Load Balancer Service



Deployment



Environment Variables



Init Containers



Config Map



Environment Variables



Volumes



Volume Mounts

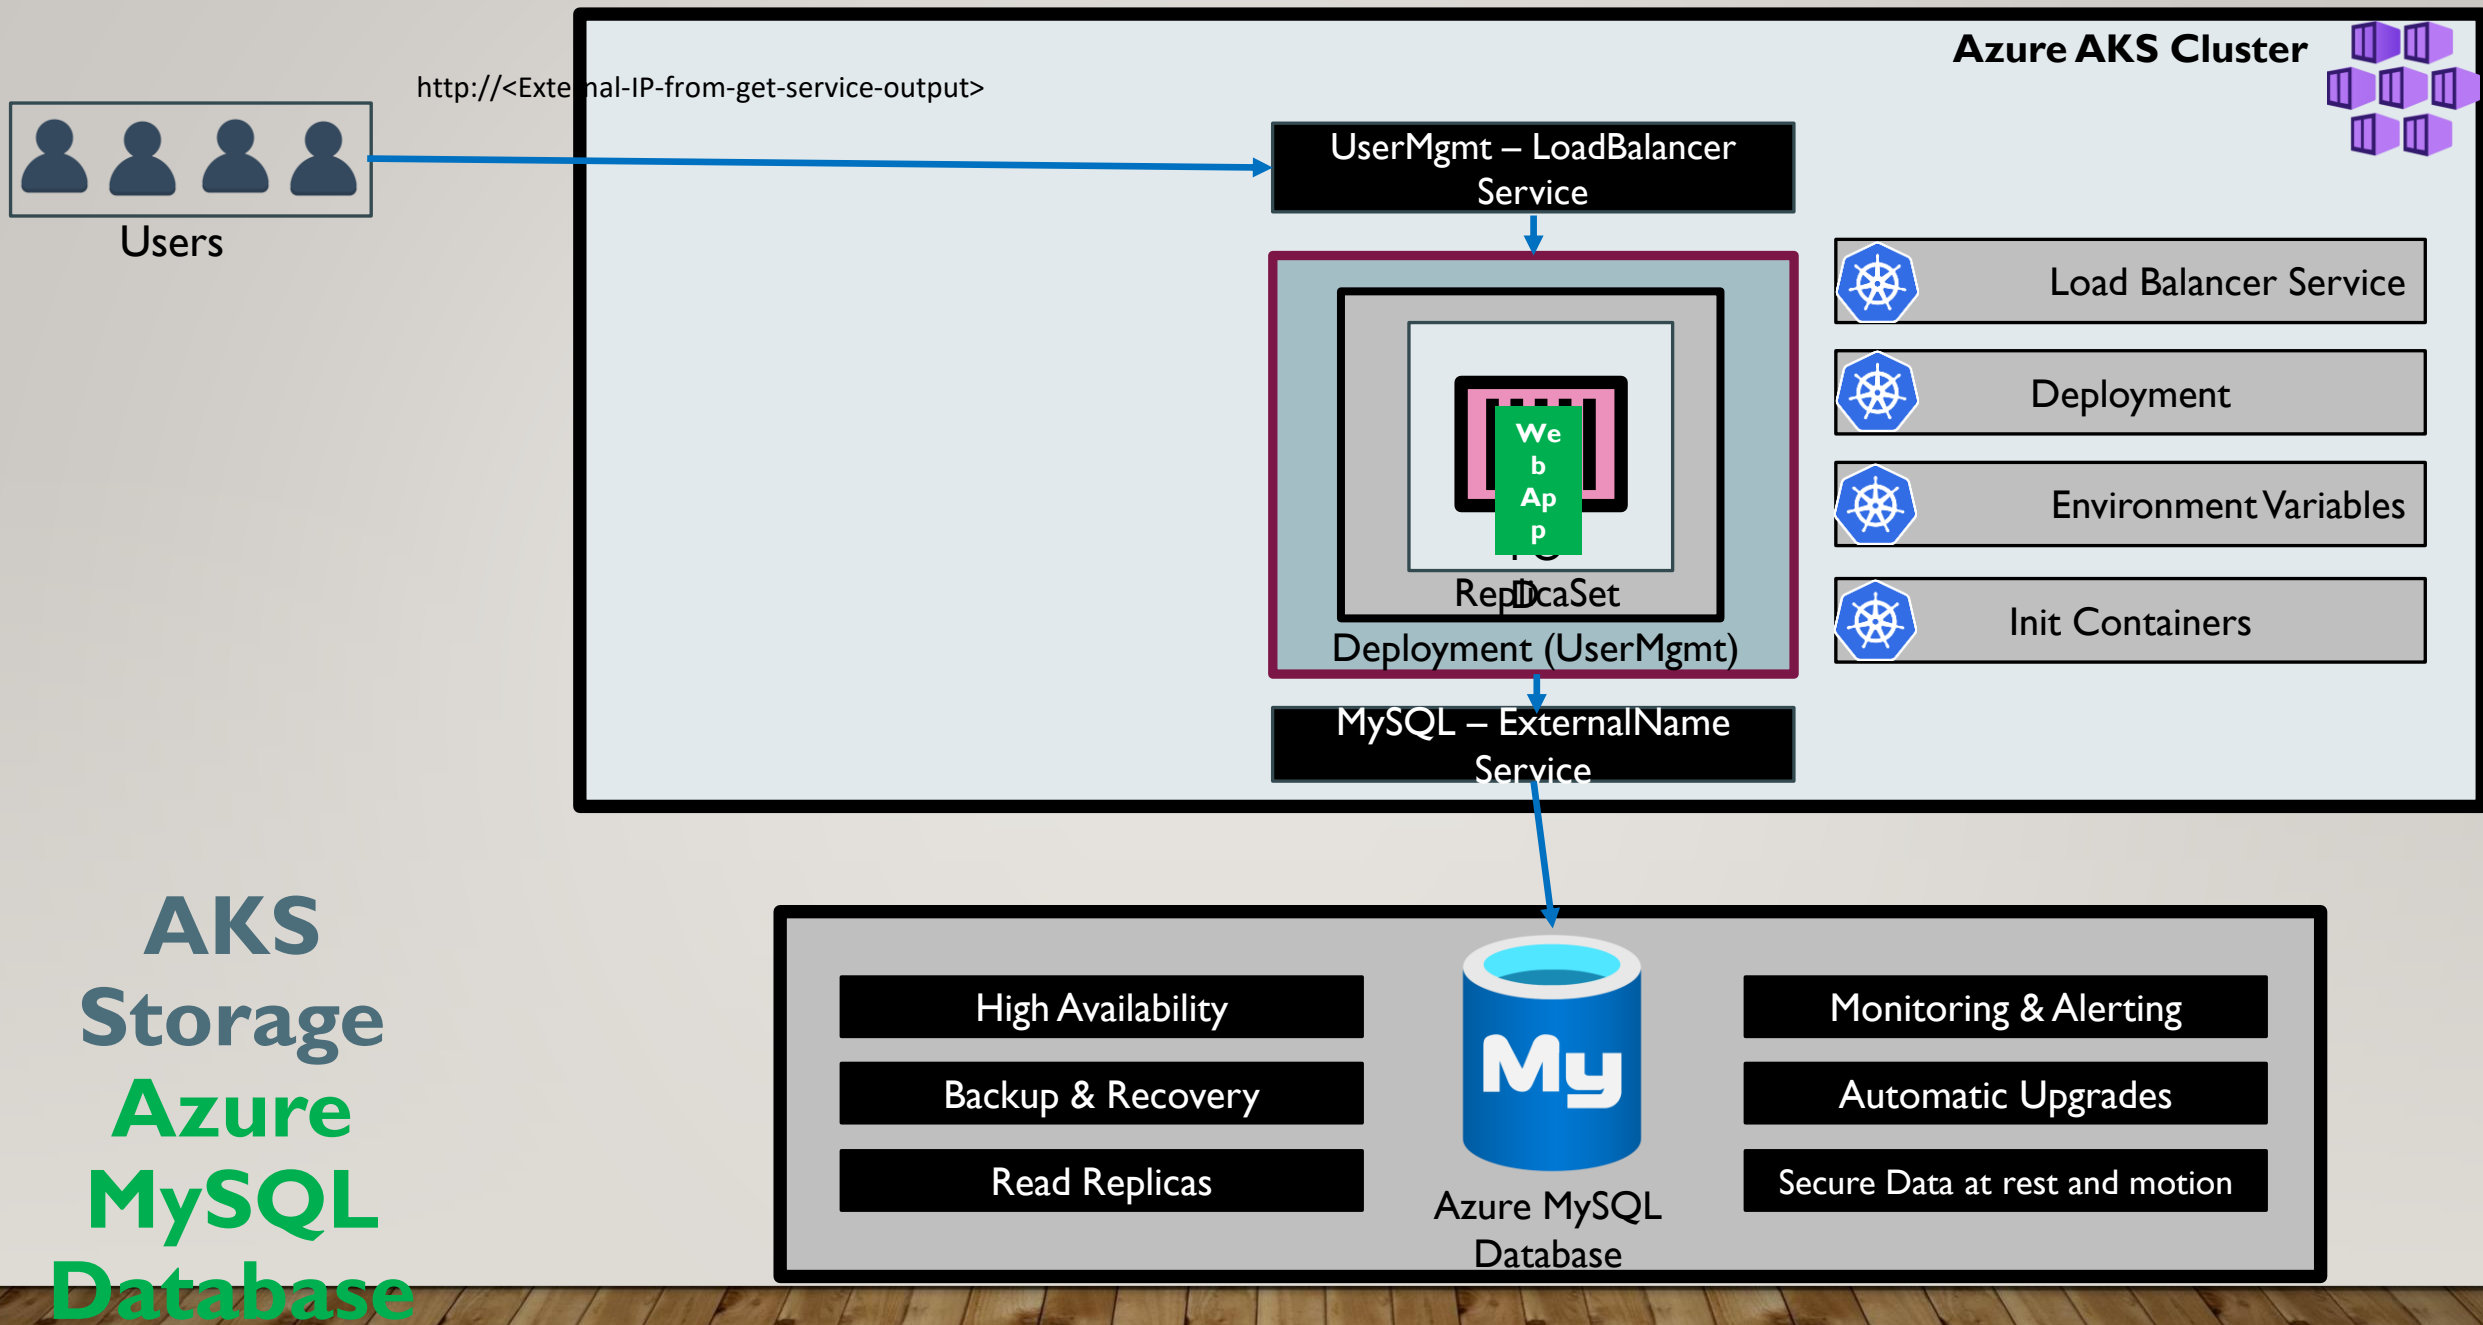


ClusterIP Service

AZURE MYSQL DATABASE



- **Features**
- Built-in **high availability** with no additional cost.
- Predictable **performance**, using inclusive pay-as-you-go pricing.
- **Scale** as needed within seconds.
- Secured to protect sensitive **data at-rest and in-motion**.
- Automatic **backups** and point-in-time-restore for up to 35 days.
- Enterprise-grade **security** and compliance.



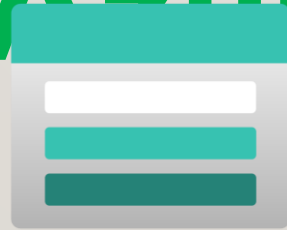


Azure AKS Storage

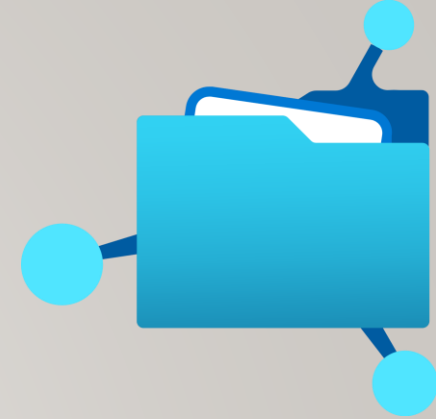
with



Azure Files



AZURE FILES



- These are simple, secure and **fully managed** cloud file shares
- We can secure **data at rest and in-transit** using SMB 3.0 and HTTPS
- We can create **high-performance file shares** using the Premium Files storage tier
- We can replace or supplement **on-premises file servers**
- **Scripting and tooling:** **PowerShell** cmdlets and **Azure CLI** can be used to create, mount, and manage Azure file shares as part of the administration of Azure applications.
- We can create and manage Azure file shares using **Azure portal** and Azure **Storage Explorer**.
- For Application workloads we can use for use cases like **Static Content storage**, shared **configuration access** to multiple JVMs etc.





Users

AKS Storage Azure Files

Key Advantage with Azure
File Shares:
Multiple pods can access
the single file share

